

Large-Scale Graph Processing

Amir H. Payberah

Swedish Institute of Computer Science

amir@sics.se


May 13-15, 2014





- ▶ **Graphs** provide a flexible abstraction for describing relationships between **discrete objects**.
- ▶ Many problems can be modeled by graphs and solved with appropriate **graph algorithms**.

Large Graph




Large-Scale Graph Processing

- ▶ Large graphs need **large-scale processing**.
- ▶ A large graph either **cannot fit into memory** of single computer or it fits with huge cost.


Question

Can we use platforms like MapReduce or Spark, which are based on **data-parallel** model, for large-scale graph proceeding?



Data-Parallel Model for Large-Scale Graph Processing

- ▶ The platforms that have worked well for developing **parallel applications** are not necessarily effective for **large-scale graph** problems.
- ▶ Why?



Graph Algorithms Characteristics (1/2)

► Unstructured problems

- Difficult to extract parallelism based on partitioning of the data: the irregular structure of graphs.
- Limited scalability: unbalanced computational loads resulting from poorly partitioned data.

Graph Algorithms Characteristics (1/2)

► Unstructured problems

- Difficult to extract parallelism based on partitioning of the data: the irregular structure of graphs.
- Limited scalability: unbalanced computational loads resulting from poorly partitioned data.

► Data-driven computations

- Difficult to express parallelism based on partitioning of computation: the structure of computations in the algorithm is not known a priori.
- The computations are dictated by nodes and links of the graph.

Graph Algorithms Characteristics (2/2)

- ▶ Poor data locality
 - The computations and data access patterns do not have much locality: the irregular structure of graphs.

Graph Algorithms Characteristics (2/2)


- ▶ Poor data locality
 - The computations and data access patterns do not have much locality: the irregular structure of graphs.
- ▶ High data access to computation ratio
 - Graph algorithms are often based on exploring the structure of a graph to perform computations on the graph data.
 - Runtime can be dominated by waiting memory fetches: low locality.

Proposed Solution

Graph-Parallel Processing

Proposed Solution


Graph-Parallel Processing



- ▶ Computation typically depends on the **neighbors**.


Graph-Parallel Processing

- ▶ Restricts the **types of computation**.
- ▶ New techniques to **partition and distribute graphs**.
- ▶ Exploit graph structure.
- ▶ Executes graph algorithms orders-of-magnitude faster than more general **data-parallel** systems.



Data-Parallel vs. Graph-Parallel Computation


Data-Parallel



Graph-Parallel




Property Graph



Data-Parallel vs. Graph-Parallel Computation

- ▶ Data-parallel computation
 - Record-centric view of data.
 - Parallelism: processing **independent** data on separate resources.

- ▶ Graph-parallel computation
 - Vertex-centric view of graphs.
 - Parallelism: partitioning graph (**dependent**) data across processing resources, and **resolving dependencies (along edges)** through **iterative** computation and communication.




Outline

- ▶ Pregel
- ▶ GraphLab
- ▶ PowerGraph
- ▶ GraphX



Seven Bridges of Königsberg

- ▶ Finding a walk through the city that would cross each bridge once and only once.
- ▶ Euler proved that the problem has no solution.




Map of Königsberg in Euler's time, highlighting the river Pregel and the bridges.

- ▶ Large-scale **graph-parallel** processing platform developed at Google.
- ▶ Inspired by **bulk synchronous parallel (BSP)** model.

Bulk Synchronous Parallel (1/2)

- ▶ It is a parallel programming model.
- ▶ The model consists of:
 - A set of processor-memory pairs.
 - A communications network that delivers messages in a point-to-point manner.
 - A mechanism for the efficient barrier synchronization for all or a subset of the processes.
 - There are no special combining, replicating, or broadcasting facilities.

Bulk Synchronous Parallel (2/2)



All vertices update in parallel (at the same time).

Vertex-Centric Programs

- ▶ Think like a vertex.
- ▶ Each vertex computes **individually** its value: in **parallel**
- ▶ Each vertex can see its **local** context, and updates its value accordingly.

- ▶ A **directed graph** that stores the program **state**, e.g., the current value.

Execution Model (1/3)

- ▶ Applications run in sequence of **iterations**: **supersteps**

Execution Model (1/3)

- ▶ Applications run in sequence of **iterations**: **supersteps**
- ▶ During a superstep, user-defined functions for each **vertex** is invoked (method **Compute()**): in **parallel**

Execution Model (1/3)

- ▶ Applications run in sequence of **iterations**: **supersteps**
- ▶ During a superstep, user-defined functions for each **vertex** is invoked (method **Compute()**): in **parallel**
- ▶ A vertex in superstep **S** can:
 - **reads** messages sent to it in superstep **S-1**.
 - **sends** messages to other vertices: receiving at superstep **S+1**.
 - **modifies** its state.

Execution Model (1/3)

- ▶ Applications run in sequence of **iterations**: **supersteps**
- ▶ During a superstep, user-defined functions for each **vertex** is invoked (method **Compute()**): in **parallel**
- ▶ A vertex in superstep **S** can:
 - **reads** messages sent to it in superstep **S-1**.
 - **sends** messages to other vertices: receiving at superstep **S+1**.
 - **modifies** its state.
- ▶ Vertices communicate directly with one another by **sending messages**.

Execution Model (2/3)

- ▶ Superstep 0: all vertices are in the active state.

Execution Model (2/3)


- ▶ Superstep 0: all vertices are in the active state.
- ▶ A vertex deactivates itself by voting to halt: no further work to do.

Execution Model (2/3)

- ▶ Superstep 0: all vertices are in the active state.
- ▶ A vertex deactivates itself by voting to halt: no further work to do.
- ▶ A halted vertex can be active if it receives a message.

Execution Model (2/3)

- ▶ Superstep 0: all vertices are in the active state.
- ▶ A vertex deactivates itself by voting to halt: no further work to do.
- ▶ A halted vertex can be active if it receives a message.
- ▶ The whole algorithm terminates when:
 - All vertices are simultaneously inactive.
 - There are no messages in transit.



Execution Model (3/3)

- ▶ **Aggregation**: a mechanism for **global** communication, monitoring, and data.

Execution Model (3/3)

- ▶ **Aggregation**: a mechanism for **global** communication, monitoring, and data.
- ▶ Runs after each **superstep**.
- ▶ Each **vertex** can provide a value to an aggregator in superstep **S**.
- ▶ The system **combines** those values and the resulting value is made available to all vertices in superstep **S + 1**.

Execution Model (3/3)


- ▶ **Aggregation**: a mechanism for **global** communication, monitoring, and data.
- ▶ Runs after each **superstep**.
- ▶ Each **vertex** can provide a value to an aggregator in superstep **S**.
- ▶ The system **combines** those values and the resulting value is made available to all vertices in superstep **S + 1**.
- ▶ A number of **predefined aggregators**, e.g., **min**, **max**, **sum**.
- ▶ Aggregation operators should be **commutative** and **associative**.

Example: Max Value (1/4)

```
i_val := val


for each message m
    if m > val then val := m

if i_val == val then
    vote_to_halt
else
    for each neighbor v
        send_message(v, val)
```



Example: Max Value (2/4)

```
i_val := val  
  
for each message m  
  if m > val then val := m  
  
if i_val == val then  
  vote_to_halt  
else  
  for each neighbor v  
    send_message(v, val)
```




Example: Max Value (3/4)

```
i_val := val


for each message m
    if m > val then val := m

if i_val == val then
    vote_to_halt
else
    for each neighbor v
        send_message(v, val)
```



Example: Max Value (4/4)


```
i_val := val  
  
for each message m  
  if m > val then val := m  
  
if i_val == val then  
  vote_to_halt  
else  
  for each neighbor v  
    send_message(v, val)
```



Example: PageRank

- ▶ Update ranks in **parallel**.
- ▶ **Iterate** until convergence.

$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$



Example: PageRank

```
Pregel_PageRank(i, messages):
    // receive all the messages
    total = 0
    foreach(msg in messages):
        total = total + msg

    // update the rank of this vertex
    R[i] = 0.15 + total

    // send new messages to neighbors
    foreach(j in out_neighbors[i]):
        sendmsg(R[i] * wij) to vertex j
```

$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

Partitioning the Graph

- ▶ The pregel library divides a graph into a number of **partitions**.
- ▶ Each consisting of a set of **vertices** and all of those vertices' **outgoing edges**.
- ▶ Vertices are assigned to partitions based on their **vertex-ID** (e.g., $\text{hash}(\text{ID})$).

Implementation (1/4)

- ▶ Master-worker model.
- ▶ User programs are copied on machines.
- ▶ One copy becomes the master.

Implementation (2/4)

- ▶ The **master** is responsible for
 - Coordinating workers activity.
 - Determining the **number of partitions**.
- ▶ Each **worker** is responsible for
 - Maintaining the **state** of its partitions.
 - Executing the user's **Compute()** method on its vertices.
 - Managing **messages** to and from other workers.

Implementation (3/4)

- ▶ The master assigns one or more **partitions** to each **worker**.

Implementation (3/4)


- ▶ The master assigns one or more partitions to each worker.
- ▶ The master assigns a portion of user input to each worker.
 - Set of records containing an arbitrary number of vertices and edges.
 - If a worker loads a vertex that belongs to that worker's partitions, the appropriate data structures are immediately updated.
 - Otherwise the worker enqueues a message to the remote peer that owns the vertex.

Implementation (4/4)

- ▶ After the **input has finished loading**, all vertices are marked as **active**.
- ▶ The master instructs each worker to perform a **superstep**.
- ▶ After the computation **halts**, the master may instruct each worker to save its portion of the graph.

Combiner

- ▶ Sending a message between workers incurs some **overhead**: use **combiner**.
- ▶ This can be reduced in some cases: sometimes vertices only care about a **summary value** for the messages it is sent (e.g., **min**, **max**, **sum**, **avg**).



Fault Tolerance (1/2)

- ▶ Fault tolerance is achieved through **checkpointing**.
- ▶ At **start of each superstep**, master tells workers to **save** their state:
 - Vertex values, edge values, incoming messages
 - Saved to persistent storage
- ▶ Master saves **aggregator values** (if any).
- ▶ This is **not** necessarily done at every superstep: **costly**

Fault Tolerance (2/2)

- ▶ When master **detects** one or more **worker failures**:
 - All workers revert to last **checkpoint**.
 - Continue **from there**.
 - That is a lot of **repeated work**.
 - At least it is better than redoing the whole job.

Pregel Summary

- ▶ Bulk Synchronous Parallel model
- ▶ Vertex-centric
- ▶ Superstep: sequence of iterations
- ▶ Master-worker model
- ▶ Communication: message passing

Pregel Limitations


- ▶ Inefficient if different regions of the graph converge at **different speed**.
- ▶ Can suffer if one **task** is **more expensive** than the others.
- ▶ Runtime of each phase is determined by the **slowest** machine.



- ▶ A directed graph that stores the program state, called data graph.


Vertex Scope

- The **scope** of vertex v is the data stored in **vertex v** , in all **adjacent vertices** and **adjacent edges**.



Execution Model (1/4)

- Rather than adopting a **message passing** as in Pregel, GraphLab allows the user defined function of a vertex to **read** and **modify** any of the data in its **scope**.



Execution Model (2/4)

- ▶ **Update** function: user-defined function similar to **Compute** in Pregel.
- ▶ Can **read** and **modify** the data within the **scope** of a vertex.
- ▶ **Schedules** the future execution of other update functions.

Execution Model (3/4)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ After executing an update function (f, g, \dots) the modified scope data in \mathcal{S}_v is written back to the data graph.
- ▶ Each task in the set of tasks \mathcal{T} , is a tuple (f, v) consisting of an update function f and a vertex v .

Execution Model (4/4)

- ▶ Sync function: similar to aggregate in Pregel.
- ▶ Maintains global aggregates.
- ▶ Performs periodically in the background.

Example: PageRank

```
GraphLab_PageRank(i)
    // compute sum over neighbors
    total = 0
    foreach(j in in_neighbors(i)):
        total = total + R[j] * wji

    // update the PageRank
    R[i] = 0.15 + total

    // trigger neighbors to run again
    foreach(j in out_neighbors(i)):
        signal vertex-program on j
```


$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

Data Consistency (1/3)

- ▶ Overlapped scopes: **race-condition** in simultaneous execution of two update functions.

Data Consistency (1/3)


- ▶ Overlapped scopes: **race-condition** in simultaneous execution of two update functions.



- ▶ **Full consistency**: during the execution $f(v)$, no other function reads or modifies data within the v scope.

Data Consistency (1/3)


- Overlapped scopes: race-condition in simultaneous execution of two update functions.



- Full consistency: during the execution $f(v)$, no other function reads or modifies data within the v scope.
- Edge consistency: during the execution $f(v)$, no other function reads or modifies any of the data on v or any of the edges adjacent to v .


Data Consistency (1/3)

- Overlapped scopes: race-condition in simultaneous execution of two update functions.



- Full consistency: during the execution $f(v)$, no other function reads or modifies data within the v scope.
- Edge consistency: during the execution $f(v)$, no other function reads or modifies any of the data on v or any of the edges adjacent to v .
- Vertex consistency: during the execution $f(v)$, no other function will be applied to v .

Data Consistency (2/3)



[Low, Y., GraphLab: A Distributed Abstraction for Large Scale Machine Learning (Doctoral dissertation, University of California), 2013.]

Data Consistency (3/3)

- ▶ Proving the **correctness** of a parallel algorithm: **sequential consistency**

Data Consistency (3/3)

- ▶ Proving the **correctness** of a parallel algorithm: **sequential consistency**
- ▶ **Sequential consistency:** if for every **parallel** execution, there exists a **sequential** execution of update functions that produces an **equivalent result**.

Data Consistency (3/3)

- ▶ Proving the **correctness** of a parallel algorithm: **sequential consistency**
- ▶ **Sequential consistency:** if for every **parallel** execution, there exists a **sequential** execution of update functions that produces an **equivalent result**.
- ▶ A simple method to **achieve serializability** is to ensure that the **scopes** of concurrently executing update functions **do not overlap**.
 - The **full consistency** model is used.
 - The **edge consistency** model is used and update functions do not modify data in adjacent vertices.
 - The **vertex consistency** model is used and update functions only access local vertex data.

GraphLab Implementation

- ▶ Shared memory implementation
- ▶ Distributed implementation

GraphLab Implementation

- ▶ Shared memory implementation
- ▶ Distributed implementation

Tasks Schedulers (1/2)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ In what **order** should the **tasks** (**vertex-update function pairs**) be called?

Tasks Schedulers (1/2)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ In what **order** should the **tasks** (**vertex-update function pairs**) be called?
 - A collection of base schedules, e.g., round-robin, and synchronous.
 - **Set scheduler:** enables users to compose custom update schedules.

Tasks Schedulers (2/2)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ How to add **new task** in the queue?

Tasks Schedulers (2/2)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ How to add **new task** in the queue?


- **FIFO:** only permits task **creation** but do **not permit task reordering**.
- **Prioritized:** **permits task reordering** at the cost of increased overhead.

Consistency

- ▶ Implemented in C++ using PThreads for parallelism.
- ▶ Consistency: **read-write lock**

Consistency

- Implemented in C++ using PThreads for parallelism.
- Consistency: **read-write lock**
- **Vertex consistency**
 - Central vertex (**write-lock**)
- **Edge consistency**
 - Central vertex (**write-lock**)
 - Adjacent vertices (**read-locks**)
- **Full consistency**
 - Central vertex (**write-locks**)
 - Adjacent vertices (**write-locks**)
- **Deadlocks** are avoided by acquiring locks sequentially following a canonical order.



GraphLab Implementation

- ▶ Shared memory implementation
- ▶ Distributed implementation

- ▶ **Graph partitioning**

- How to efficiently load, partition and distribute the data graph across machines?


- ▶ **Consistency**

- How to achieve consistency in the distributed setting?


- ▶ **Fault tolerance**

Graph Partitioning - Phase 1 (1/2)

- ▶ Two-phase partitioning.
- ▶ Partitioning the data graph into k parts, called **atom**: $k \gg$ number of machines.
- ▶ **meta-graph**: the graph of atoms (one vertex for each atom).
- ▶ **Atom weight**: the amount of data it stores.
- ▶ **Edge weight**: the number of edges crossing the atoms.



Graph Partitioning - Phase 1 (2/2)



- ▶ Each **atom** is stored as a separate file on a distributed storage system, e.g., HDFS.
- ▶ Each atom file is a simple **binary** that stores **interior** and the **ghosts** of the partition information.
- ▶ **Ghost**: set of **vertices and edges adjacent** to the partition boundary.

Graph Partitioning - Phase 2

- ▶ Meta-graph is very **small**.
- ▶ A **fast balanced partition** of the **meta-graph** over the physical machines.
- ▶ Assigning graph atoms to machines.

- ▶ To achieve a **serializable** parallel execution of a set of **dependent tasks**.
- ▶ Chromatic Engine
- ▶ Distributed Locking Engine

Consistency - Chromatic Engine

- ▶ Construct a **vertex coloring**: assigns a color to each vertex such that **no adjacent** vertices share the same color.

Consistency - Chromatic Engine

- ▶ Construct a **vertex coloring**: assigns a color to each vertex such that **no adjacent** vertices share the same color.
- ▶ **Edge consistency**: executing, synchronously, all update tasks associated with vertices of the **same** color before proceeding to the next color.

Consistency - Chromatic Engine

- ▶ Construct a **vertex coloring**: assigns a color to each vertex such that **no adjacent** vertices share the same color.
- ▶ **Edge consistency**: executing, synchronously, all update tasks associated with vertices of the **same** color before proceeding to the next color.
- ▶ **Full consistency**: no vertex shares the same color as any of its distance **two** neighbors.

Consistency - Chromatic Engine

- ▶ Construct a **vertex coloring**: assigns a color to each vertex such that **no adjacent** vertices share the same color.
- ▶ **Edge consistency**: executing, synchronously, all update tasks associated with vertices of the **same** color before proceeding to the next color.
- ▶ **Full consistency**: no vertex shares the same color as any of its distance **two** neighbors.
- ▶ **Vertex consistency**: assigning all vertices the **same** color.

Consistency - Distributed Locking Engine

- ▶ Associating a **readers-writer** lock with each vertex.
- ▶ **Vertex consistency**
 - Central vertex (**write-lock**)
- ▶ **Edge consistency**
 - Central vertex (**write-lock**), Adjacent vertices (**read-locks**)
- ▶ **Full consistency**
 - Central vertex (**write-locks**), Adjacent vertices (**write-locks**)
- ▶ **Deadlocks** are avoided by acquiring locks sequentially following a canonical order.

Fault Tolerance - Synchronous

- ▶ The systems **periodically** signals all computation activity to **halt**.
- ▶ Then synchronizes all caches (ghosts) and **saves to disk** all data which has been modified since the last snapshot.
- ▶ **Simple**, but eliminates the systems advantage of **asynchronous** computation.

Fault Tolerance - Asynchronous

- ▶ Based on the Chandy-Lamport algorithm.
- ▶ The **snapshot** function is implemented as an **update function** in vertices.
- ▶ The Snapshot update takes **priority** over all other update functions.
- ▶ **Edge Consistency** is used on all update functions.

```
if v was already snapshotted then
    ↘ Quit
    Save  $D_v$  // Save current vertex
    // Save all edges connected to un-snapshotted vertices
    foreach  $u \in N[v]$  do                                // Loop over neighbors
        if u was not snapshotted then
            ↗ Save  $D_{u \rightarrow v}$  if edge  $u \rightarrow v$  exists
            ↗ Save  $D_{v \rightarrow u}$  if edge  $v \rightarrow u$  exists
            ↗ Reschedule u for a Snapshot Update
    ↗ Mark v as snapshotted
```


- ▶ Asynchronous model
- ▶ Vertex-centric
- ▶ Communication: distributed shared memory
- ▶ Three consistency levels: full, edge-level, and vertex-level

GraphLab Limitations

- ▶ Poor performance on **Natural** graphs.


Natural Graphs

- ▶ Graphs derived from **natural phenomena**.
- ▶ Skewed **Power-Law** degree distribution.



Natural Graphs Challenges

- ▶ Traditional graph-partitioning algorithms (**edge-cut** algorithms) perform **poorly** on Power-Law Graphs.
- ▶ Challenges of **high-degree** vertices.




Proposed Solution


Vertex-Cut Partitioning

Proposed Solution

Vertex-Cut Partitioning







Edge-cut




Vertex-cut


Edge-cut vs. Vertex-cut Partitioning



Edge-cut vs. Vertex-cut Partitioning



Edge-cut



Vertex-cut


PowerGraph

- ▶ Vertex-cut partitioning of graphs.
- ▶ Factorizes the GraphLab **update function** into the **Gather**, **Apply** and **Scatter** phases (GAS).

Gather-Apply-Scatter Programming Model

► Gather


- Accumulate information about neighborhood through a generalized sum.



Gather-Apply-Scatter Programming Model

► Gather

- Accumulate information about neighborhood through a generalized sum.




► Apply

- Apply the accumulated value to center vertex.

Gather-Apply-Scatter Programming Model

► Gather

- Accumulate information about neighborhood through a generalized sum.




► Apply

- Apply the accumulated value to center vertex.

► Scatter

- Update adjacent edges and vertices.



- ▶ A directed graph that stores the program state, called data graph.

Execution Model (1/2)

- ▶ Vertex-centric programming: implementing the `GASVertexProgram` interface (`vertex-program` for short).
- ▶ Similar to `Computation` in `Pregel`, and `update` function in `GraphLab`.

```
interface GASVertexProgram(u) {
    // Run on gather_nbrs(u)
    gather(Du, Du-v, Dv) → Accum
    sum(Accum left, Accum right) → Accum
    apply(Du, Accum) → Dunew
    // Run on scatter_nbrs(u)
    scatter(Dunew, Du-v, Dv) → (Du-vnew, Accum)
}
```

Execution Model (2/2)

```
Input: Center vertex u
if Cache Disabled then
    // Basic Gather-Apply-Scatter Model
    foreach neighbor v in gather_nbrs(u) do
        au ← sum(au, gather(Du, Du-v, Dv))
        Du ← apply(Du, au)
        foreach neighbor v scatter_nbrs(u) do
            (Du-v) ← scatter(Du, Du-v, Dv)
else if Cache Enabled then
    // Faster GAS Model with Delta Caching
    if cached accumulator au is empty then
        foreach neighbor v in gather_nbrs(u) do
            au ← sum(au, gather(Du, Du-v, Dv))
            Du ← apply(Du, au)
            foreach neighbor v scatter_nbrs(u) do
                (Du-v, Δa) ← scatter(Du, Du-v, Dv)
                if av and Δa are not Empty then av ← sum(av, Δa)
                else av ← Empty
```

Execution Model (2/2)

Input: Center vertex u

if Cache Disabled then

```
// Basic Gather-Apply-Scatter Model
foreach neighbor  $v$  in  $\text{gather\_nbrs}(u)$  do
     $a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$ 
 $D_u \leftarrow \text{apply}(D_u, a_u)$ 
foreach neighbor  $v$  in  $\text{scatter\_nbrs}(u)$  do
     $(D_{u-v}) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$ 
```

else if Cache Enabled then

```
// Faster GAS Model with Delta Caching
if cached accumulator  $a_u$  is empty then
    foreach neighbor  $v$  in  $\text{gather\_nbrs}(u)$  do
         $a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$ 
 $D_u \leftarrow \text{apply}(D_u, a_u)$ 
foreach neighbor  $v$  in  $\text{scatter\_nbrs}(u)$  do
     $(D_{u-v}, \Delta a) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$ 
    if  $a_v$  and  $\Delta a$  are not Empty then  $a_v \leftarrow \text{sum}(a_v, \Delta a)$ 
    else  $a_v \leftarrow \text{Empty}$ 
```

Example: PageRank

```
PowerGraph_PageRank(i):
    Gather(j -> i):
        return wji * R[j]

    sum(a, b):
        return a + b

    // total: Gather and sum
    Apply(i, total):
        R[i] = 0.15 + total

    Scatter(i -> j):
        if R[i] changed then activate(j)
```

$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

Scheduling (1/5)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ PowerGraph inherits the **dynamic scheduling** of **GraphLab**.

Scheduling (2/5)

- ▶ Initially all vertices are active.

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty **do**
1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$
Output: Modified Data Graph $G' = (V, E, D')$

Scheduling (2/5)

- ▶ Initially all vertices are active.

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty do

1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G' = (V, E, D')$

- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.

Scheduling (2/5)

- ▶ Initially all vertices are active.

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty do

1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G' = (V, E, D')$

- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.
- ▶ The order of executing activated vertices is up to the PowerGraph execution engine.

Scheduling (2/5)

- ▶ Initially all vertices are active.

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty do

1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G' = (V, E, D')$

- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.
- ▶ The order of executing activated vertices is up to the PowerGraph execution engine.
- ▶ Once a vertex-program completes the scatter phase it becomes inactive until it is reactivated.

Scheduling (2/5)

- ▶ Initially all vertices are active.

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty do

1 $(f, v) \leftarrow \text{RemoveNext } (\mathcal{T})$
2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$

Output: Modified Data Graph $G = (V, E, D')$

- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.
- ▶ The order of executing activated vertices is up to the PowerGraph execution engine.
- ▶ Once a vertex-program completes the scatter phase it becomes inactive until it is reactivated.
- ▶ Vertices can activate themselves and neighboring vertices.

Scheduling (3/5)

- ▶ PowerGraph can execute both **synchronously** and **asynchronously**.
 - Bulk synchronous execution
 - Asynchronous execution

Scheduling - Bulk Synchronous Execution (4/5)

- ▶ Similar to Pregel.

- ▶ Similar to Pregel.
- ▶ Minor-step: executing the `gather`, `apply`, and `scatter` in order.
 - Runs `synchronously` on all `active` vertices with a `barrier` at the end.

- ▶ Similar to Pregel.
- ▶ Minor-step: executing the gather, apply, and scatter **in order**.
 - Runs **synchronously** on all **active** vertices with a **barrier** at the end.
- ▶ Super-step: a complete series of GAS minor-steps.

Scheduling - Bulk Synchronous Execution (4/5)

- ▶ Similar to Pregel.
- ▶ Minor-step: executing the **gather**, **apply**, and **scatter** **in order**.
 - Runs **synchronously** on all **active** vertices with a **barrier** at the end.
- ▶ Super-step: a complete series of GAS minor-steps.
- ▶ Changes made to the vertex/edge data are committed at the **end** of each **minor-step** and are visible in the **subsequent minor-steps**.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.
 - **Dining philosophers** problem, where each vertex is a philosopher, and each edge is a fork.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.
 - **Dining philosophers** problem, where each vertex is a philosopher, and each edge is a fork.
 - **GraphLab** implements **Dijkstras** solution, where forks are acquired **sequentially** according to a total ordering.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.
 - **Dining philosophers** problem, where each vertex is a philosopher, and each edge is a fork.
 - **GraphLab** implements **Dijkstras** solution, where forks are acquired **sequentially** according to a total ordering.
 - **PowerGraph** implements **Chandy-Misra** solution, which acquires all forks **simultaneously**.

Delta Caching (1/2)

- ▶ Changes in a few of its neighbors → triggering a vertex-program
- ▶ The gather operation is invoked on all neighbors: wasting computation cycles

Delta Caching (1/2)

- ▶ Changes in a few of its neighbors → triggering a vertex-program
- ▶ The gather operation is invoked on all neighbors: wasting computation cycles
- ▶ Maintaining a cache of the accumulator a_v from the previous gather phase for each vertex.
- ▶ The scatter can return an additional Δa , which is added to the cached accumulator a_v .

Delta Caching (2/2)

Input: Center vertex u

if Cache Disabled then

```
// Basic Gather-Apply-Scatter Model
foreach neighbor  $v$  in  $\text{gather\_nbrs}(u)$  do
     $a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$ 
 $D_u \leftarrow \text{apply}(D_u, a_u)$ 
foreach neighbor  $v$  in  $\text{scatter\_nbrs}(u)$  do
     $(D_{u-v}) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$ 
```

else if Cache Enabled then

~~```
// Faster GAS Model with Delta Caching
if cached accumulator a_u is empty then
 foreach neighbor v in $\text{gather_nbrs}(u)$ do
 $a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$
 $D_u \leftarrow \text{apply}(D_u, a_u)$
foreach neighbor v in $\text{scatter_nbrs}(u)$ do
 $(D_{u-v}, \Delta a) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$
 if a_v and Δa are not Empty then $a_v \leftarrow \text{sum}(a_v, \Delta a)$
 else $a_v \leftarrow \text{Empty}$
```~~

## Delta Caching (2/2)

```
Input: Center vertex u
if Cache Disabled then
 // Basic Gather-Apply-Scatter Model
 foreach neighbor v in gather_nbrs(u) do
 au ← sum(au, gather(Du, Du-v, Dv))
 Du ← apply(Du, au)
 foreach neighbor v scatter_nbrs(u) do
 (Du-v) ← scatter(Du, Du-v, Dv)
else if Cache Enabled then
 // Faster GAS Model with Delta Caching
 if cached accumulator au is empty then
 foreach neighbor v in gather_nbrs(u) do
 au ← sum(au, gather(Du, Du-v, Dv))
 Du ← apply(Du, au)
 foreach neighbor v scatter_nbrs(u) do
 (Du-v, Δa) ← scatter(Du, Du-v, Dv)
 if av and Δa are not Empty then av ← sum(av, Δa)
 else av ← Empty
```

## Example: PageRank (Delta-Caching)

```
PowerGraph_PageRank(i):
 Gather(j -> i):
 return wji * R[j]

 sum(a, b):
 return a + b

 // total: Gather and sum
 Apply(i, total):
 new = 0.15 + total
 R[i].delta = new - R[i]
 R[i] = new

 Scatter(i -> j):
 if R[i] changed then activate(j)
 return wij * R[i].delta
```

$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

# Graph Partitioning

- ▶ Vertex-cut partitioning.
- ▶ Evenly assign edges to machines.
  - Minimize machines spanned by each vertex.
- ▶ Two proposed solutions:
  - Random edge placement.
  - Greedy edge placement.

- ▶ Randomly assign edges to machines.
- ▶ Completely parallel and easy to **distribute**.
- ▶ High replication factor.

## Greedy Vertex-Cuts (1/2)

- ▶  $A(v)$ : set of machines that contain adjacent edges of  $v$ .

## Greedy Vertex-Cuts (1/2)

- ▶  $A(v)$ : set of machines that contain adjacent edges of  $v$ .
- ▶ **Case 1:** If  $A(u)$  and  $A(v)$  intersect, then the edge should be assigned to a machine in the intersection.

## Greedy Vertex-Cuts (1/2)

- ▶  $A(v)$ : set of machines that contain adjacent edges of  $v$ .
- ▶ **Case 1:** If  $A(u)$  and  $A(v)$  intersect, then the edge should be assigned to a machine in the intersection.
- ▶ **Case 2:** If  $A(u)$  and  $A(v)$  are not empty and do not intersect, then the edge should be assigned to one of the machines from the vertex with the most unassigned edges.

## Greedy Vertex-Cuts (1/2)

- ▶  $A(v)$ : set of machines that contain adjacent edges of  $v$ .
- ▶ **Case 1:** If  $A(u)$  and  $A(v)$  intersect, then the edge should be assigned to a machine in the intersection.
- ▶ **Case 2:** If  $A(u)$  and  $A(v)$  are not empty and do not intersect, then the edge should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ **Case 3:** If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.

## Greedy Vertex-Cuts (1/2)

- ▶  $A(v)$ : set of machines that contain adjacent edges of  $v$ .
- ▶ **Case 1:** If  $A(u)$  and  $A(v)$  intersect, then the edge should be assigned to a machine in the intersection.
- ▶ **Case 2:** If  $A(u)$  and  $A(v)$  are not empty and do not intersect, then the edge should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ **Case 3:** If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.
- ▶ **Case 4:** If neither vertex has been assigned, then assign the edge to the least loaded machine.

## Greedy Vertex-Cuts (2/2)

- ▶ **Coordinated** edge placement:
  - Requires coordination to place each edge
  - Slower, but **higher** quality cuts
- ▶ **Oblivious** edge placement:
  - Approx. greedy objective without coordination
  - Faster, but **lower** quality cuts

- ▶ Gather-Apply-Scatter programming model
- ▶ Synchronous and Asynchronous models
- ▶ Vertex-cut graph partitioning


- ▶ Any limitations?

# Data-Parallel vs. Graph-Parallel Computation


- ▶ Graph-parallel computation: **restricting** the types of computation to achieve **performance**.

# Data-Parallel vs. Graph-Parallel Computation

- ▶ Graph-parallel computation: **restricting** the types of computation to achieve **performance**.
- ▶ **But**, the same restrictions make it **difficult** and **inefficient** to express many stages in a typical graph-analytics **pipeline**.




# Data-Parallel and Graph-Parallel Pipeline






- ▶ Moving between **table** and **graph** views of the **same physical data**.
- ▶ **Inefficient:** extensive **data movement** and **duplication** across the network and file system.

# GraphX vs. Data-Parallel/Graph-Parallel Systems



# GraphX vs. Data-Parallel/Graph-Parallel Systems






- ▶ New API that blurs the distinction between Tables and Graphs.
- ▶ New system that unifies Data-Parallel and Graph-Parallel systems.
- ▶ It is implemented on top of Spark.

# Unifying Data-Parallel and Graph-Parallel Analytics


- ▶ Tables and Graphs are composable views of the same physical data.
- ▶ Each view has its own operators that exploit the semantics of the view to achieve efficient execution.



# Data Model

- ▶ **Property Graph:** represented using two Spark RDDs:
  - Edge collection: VertexRDD
  - Vertex collection: EdgeRDD

```
// VD: the type of the vertex attribute
// ED: the type of the edge attribute
class Graph[VD, ED] {
 val vertices: VertexRDD[VD]
 val edges: EdgeRDD[ED]
}
```




# Primitive Data Types

```
// Vertex collection
class VertexRDD[VD] extends RDD[(VertexId, VD)]


// Edge collection
class EdgeRDD[ED] extends RDD[Edge[ED]]
case class Edge[ED](srcId: VertexId = 0, dstId: VertexId = 0,
 attr: ED = null.asInstanceOf[ED])

// Edge Triple
class EdgeTriplet[VD, ED] extends Edge[ED]
```

- ▶ **EdgeTriplet** represents an **edge** along with the **vertex attributes** of its **neighboring** vertices.



# Example (1/3)



Vertex Table

| Id | Property (V)          |
|----|-----------------------|
| 3  | (rxin, student)       |
| 7  | (jgonzal, postdoc)    |
| 5  | (franklin, professor) |
| 2  | (istoica, professor)  |

Edge Table

| SrcId | DstId | Property (E) |
|-------|-------|--------------|
| 3     | 7     | Collaborator |
| 5     | 3     | Advisor      |
| 2     | 5     | Colleague    |
| 5     | 7     | PI           |

## Example (2/3)

```
val sc: SparkContext

// Create an RDD for the vertices
val users: VertexRDD[(String, String)] = sc.parallelize(
 Array((3L, ("rxin", "student")), (7L, ("jgonzal", "postdoc")),
 (5L, ("franklin", "prof")), (2L, ("istoica", "prof"))))

// Create an RDD for edges
val relationships: EdgeRDD[String] = sc.parallelize(
 Array(Edge(3L, 7L, "collab"), Edge(5L, 3L, "advisor"),
 Edge(2L, 5L, "colleague"), Edge(5L, 7L, "pi")))

// Define a default user in case there are relationship with missing user
val defaultUser = ("John Doe", "Missing")

// Build the initial Graph
val userGraph: Graph[(String, String), String] =
 Graph(users, relationships, defaultUser)
```

## Example (3/3)

```
// Constructed from above
val userGraph: Graph[(String, String), String]

// Count all users which are postdocs
userGraph.vertices.filter((id, (name, pos)) => pos == "postdoc").count

// Count all the edges where src > dst
userGraph.edges.filter(e => e.srcId > e.dstId).count

// Use the triplets view to create an RDD of facts
val facts: RDD[String] = graph.triplets.map(triplet =>
 triplet.srcAttr._1 + " is the " +
 triplet.attr + " of " + triplet.dstAttr._1)

// Remove missing vertices as well as the edges to connected to them
val validGraph = graph.subgraph(vpred = (id, attr) => attr._2 != "Missing")

facts.collect.foreach(println(_))
```

## Property Operators (1/2)

```
class Graph[VD, ED] {
 def mapVertices[VD2](map: (VertexId, VD) => VD2): Graph[VD2, ED]

 def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]

 def mapTriplets[ED2](map: EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]
}
```

- ▶ They yield **new graphs** with the vertex or edge properties modified by the map function.
- ▶ The graph **structure** is **unaffected**.

## Property Operators (2/2)

```
val newGraph = graph.mapVertices((id, attr) => mapUdf(id, attr))
```

```
val newVertices = graph.vertices.map((id, attr) => (id, mapUdf(id, attr)))
val newGraph = Graph(newVertices, graph.edges)
```


- ▶ Both are logically equivalent, but the second one **does not preserve** the structural indices and would not benefit from the GraphX system **optimizations**.


# Map Reduce Triplets

- ▶ Map-Reduce for each vertex

mapF( )  $\Rightarrow$  

mapF( )  $\Rightarrow$  

reduceF( )  $\Rightarrow$  




# Map Reduce Triplets

- ▶ Map-Reduce for each vertex

mapF( ) → 

mapF( ) → 

reduceF( ) → 



```
// what is the age of the oldest follower for each user?
val oldestFollowerAge = graph.mapReduceTriplets(
 e => (e.dstAttr, e.srcAttr), // Map
 (a, b) => max(a, b) // Reduce
) .vertices
```

# Structural Operators

```
class Graph[VD, ED] {
 // returns a new graph with all the edge directions reversed
 def reverse: Graph[VD, ED]

 // returns the graph containing only the vertices and edges that satisfy
 // the vertex predicate
 def subgraph(epred: EdgeTriplet[VD,ED] => Boolean,
 vpred: (VertexId, VD) => Boolean): Graph[VD, ED]

 // a subgraph by returning a graph that contains the vertices and edges
 // that are also found in the input graph
 def mask[VD2, ED2](other: Graph[VD2, ED2]): Graph[VD, ED]
}
```

# Structural Operators Example

```
// Build the initial Graph
val graph = Graph(users, relationships, defaultUser)

// Run Connected Components
val ccGraph = graph.connectedComponents()

// Remove missing vertices as well as the edges to connected to them
val validGraph = graph.subgraph(vpred = (id, attr) => attr._2 != "Missing")

// Restrict the answer to the valid subgraph
val validCCGraph = ccGraph.mask(validGraph)
```

# Join Operators

- ▶ To join data from external collections (RDDs) with graphs.

```
class Graph[VD, ED] {
 // joins the vertices with the input RDD and returns a new graph
 // by applying the map function to the result of the joined vertices
 def joinVertices[U](table: RDD[(VertexId, U)])
 (map: (VertexId, VD, U) => VD): Graph[VD, ED]

 // similarly to joinVertices, but the map function is applied to
 // all vertices and can change the vertex property type
 def outerJoinVertices[U, VD2](table: RDD[(VertexId, U)])
 (map: (VertexId, VD, Option[U]) => VD2): Graph[VD2, ED]
}
```

# Graph Builders


```
// load a graph from a list of edges on disk
object GraphLoader {
 def edgeListFile(
 sc: SparkContext,
 path: String,
 canonicalOrientation: Boolean = false,
 minEdgePartitions: Int = 1)
 : Graph[Int, Int]
}

// graph file
This is a comment
2 1
4 1
1 2
```

- ▶ GraphX is implemented on top of **Spark**
- ▶ **In-memory** caching
- ▶ **Lineage-based** fault tolerance
- ▶ Programmable partitioning

# Distributed Graph Representation (1/2)

- ▶ Representing graphs using two RDDs: **edge-collection** and **vertex-collection**
- ▶ Vertex-cut partitioning (like **PowerGraph**)



## Distributed Graph Representation (2/2)

- ▶ Each vertex partition contains a **bitmask** and **routing table**.
- ▶ **Routing table:** a **logical map** from a vertex id to the set of edge partitions that contains adjacent edges.
- ▶ **Bitmask:** enables the set intersection and filtering.
  - Vertices bitmasks are updated after each operation (e.g., `mapReduceTriplets`).
  - Vertices hidden by the bitmask **do not** participate in the graph operations.

# Summary

## ► Pregel

- Synchronous model: super-step
- Message passing

## ► GraphLab

- Asynchronous model: distributed shared-memory
- Edge-cut partitioning

## ► PowerGraph

- GAS programming model
- Vertex-cut partitioning

## ► GraphX

- Unifying data-parallel and graph-parallel analytics
- Vertex-cut partitioning

# Questions?

## Acknowledgements

Some pictures were derived from the Spark web site  
(<http://spark.apache.org/>).