GraphLab: A New Framework For Parallel Machine Learning

Amir H. Payberah amir@sics.se

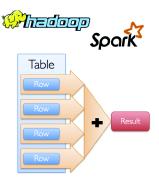
Amirkabir University of Technology (Tehran Polytechnic)



Reminder

Data-Parallel Model for Large-Scale Graph Processing

► The platforms that have worked well for developing parallel applications are not necessarily effective for large-scale graph problems.



Graph-Parallel Processing

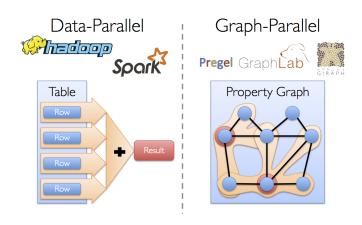
- Restricts the types of computation.
- ► New techniques to partition and distribute graphs.
- Exploit graph structure.
- Executes graph algorithms orders-of-magnitude faster than more general data-parallel systems.







Data-Parallel vs. Graph-Parallel Computation



► Vertex-centric

- ► Vertex-centric
- ► Bulk Synchronous Parallel (BSP)

- ▶ Vertex-centric
- ► Bulk Synchronous Parallel (BSP)
- ► Runs in sequence of iterations (supersteps)

- Vertex-centric
- ► Bulk Synchronous Parallel (BSP)
- Runs in sequence of iterations (supersteps)
- ► A vertex in superstep S can:
 - reads messages sent to it in superstep S-1.
 - sends messages to other vertices: receiving at superstep S+1.
 - · modifies its state.

Pregel Limitations

- ▶ Inefficient if different regions of the graph converge at different speed.
- ► Can suffer if one task is more expensive than the others.
- ▶ Runtime of each phase is determined by the slowest machine.

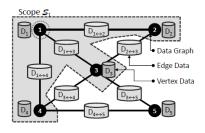


Data Model

► A directed graph that stores the program state, called data graph.

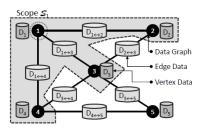
Vertex Scope

► The scope of vertex v is the data stored in vertex v, in all adjacent vertices and adjacent edges.



Programming Model (1/3)

▶ Rather than adopting a message passing as in Pregel, GraphLab allows the user defined function of a vertex to read and modify any of the data in its scope.



Programming Model (2/3)

- ▶ Update function: user-defined function similar to Compute in Pregel.
- ► Can read and modify the data within the scope of a vertex.
- ▶ Schedules the future execution of other update functions.

Programming Model (3/3)

- ► Sync function: similar to aggregate in Pregel.
- ► Maintains global aggregates.
- ▶ Performs periodically in the background.

Execution Model

```
Input: Data Graph G = (V, E, D)

Input: Initial task set \mathcal{T} = \{(f, v_1), (g, v_2), ...\}

while \mathcal{T} is not Empty do

1 (f, v) \leftarrow \text{RemoveNext}(\mathcal{T})

2 (\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)

3 \mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'

Output: Modified Data Graph G = (V, E, D')
```

Execution Model

```
Input: Data Graph G = (V, E, D)

Input: Initial task set \mathcal{T} = \{(f, v_1), (g, v_2), ...\}

while \mathcal{T} is not Empty do

1 (f, v) \leftarrow \text{RemoveNext}(\mathcal{T})

2 (\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)

3 \mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'

Output: Modified Data Graph G = (V, E, D')
```

▶ Each task in the set of tasks \mathcal{T} , is a tuple (f, v) consisting of an update function f and a vertex v.

Execution Model

```
Input: Data Graph G = (V, E, D)

Input: Initial task set \mathcal{T} = \{(f, v_1), (g, v_2), ...\}

while \mathcal{T} is not Empty do

1 (f, v) \leftarrow \text{RemoveNext}(\mathcal{T})

2 (\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)

3 \mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'

Output: Modified Data Graph G = (V, E, D')
```

- ▶ Each task in the set of tasks \mathcal{T} , is a tuple (f, v) consisting of an update function f and a vertex v.
- After executing an update function (f, g, \cdots) the modified scope data in S_v is written back to the data graph.

Example: PageRank

```
GraphLab_PageRank(i)
  // compute sum over neighbors
  total = 0
  foreach(j in in_neighbors(i)):
    total = total + R[j] * wji

// update the PageRank
  R[i] = 0.15 + total

// trigger neighbors to run again
  foreach(j in out_neighbors(i)):
    signal vertex-program on j
```

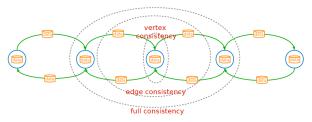
$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

Data Consistency (1/3)

 Overlapped scopes: race-condition in simultaneous execution of two update functions.

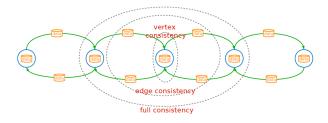
Data Consistency (1/3)

Overlapped scopes: race-condition in simultaneous execution of two update functions.



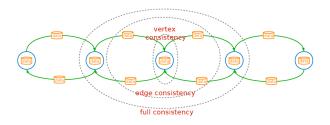
▶ Full consistency: during the execution f(v), no other function reads or modifies data within the v scope.

Data Consistency (2/3)



▶ Edge consistency: during the execution f(v), no other function reads or modifies any of the data on v or any of the edges adjacent to v.

Data Consistency (3/3)



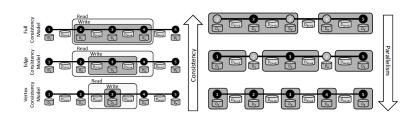
ightharpoonup Vertex consistency: during the execution f(v), no other function will be applied to v.

▶ Proving the correctness of a parallel algorithm: sequential consistency

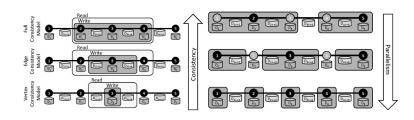
- ▶ Proving the correctness of a parallel algorithm: sequential consistency
- ► Sequential consistency: if for every parallel execution, there exists a sequential execution of update functions that produces an equivalent result.

► A simple method to achieve serializability is to ensure that the scopes of concurrently executing update functions do not overlap.

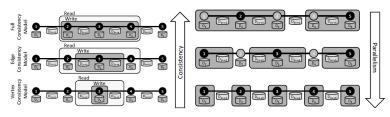
- ► A simple method to achieve serializability is to ensure that the scopes of concurrently executing update functions do not overlap.
 - The full consistency model is used.



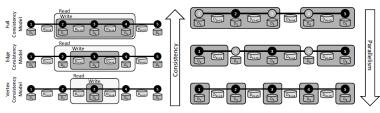
- ► A simple method to achieve serializability is to ensure that the scopes of concurrently executing update functions do not overlap.
 - The full consistency model is used.
 - The edge consistency model is used and update functions do not modify data in adjacent vertices.



- ► A simple method to achieve serializability is to ensure that the scopes of concurrently executing update functions do not overlap.
 - The full consistency model is used.
 - The edge consistency model is used and update functions do not modify data in adjacent vertices.
 - The vertex consistency model is used and update functions only access local vertex data.



Consistency vs. Parallelism



Consistency vs. Parallelism

[Low, Y., GraphLab: A Distributed Abstraction for Large Scale Machine Learning (Doctoral dissertation, University of California), 2013.]

GraphLab Implementation

- ► Shared memory implementation
- ► Distributed implementation

GraphLab Implementation

- ► Shared memory implementation
- ► Distributed implementation

Tasks Schedulers (1/2)

```
 \begin{aligned} & \textbf{Input:} \  \, \textbf{Data Graph} \ G = (V, E, D) \\ & \textbf{Input:} \  \, \textbf{Initial task set} \ \mathcal{T} = \{(f, v_1), (g, v_2), \ldots\} \\ & \textbf{while} \ \mathcal{T} \  \, \textbf{is not Empty do} \\ & \textbf{1} \\ & \textbf{2} \\ & \textbf{3} \\ & \textbf{T} \leftarrow \mathcal{T} \cup \mathcal{T}' \\ & \textbf{Output:} \  \, \textbf{Modified Data Graph} \  \, G = (V, E, D') \end{aligned}
```

► In what order should the tasks (vertex-update function pairs) be called?

Tasks Schedulers (1/2)

- ▶ In what order should the tasks (vertex-update function pairs) be called?
 - A collection of base schedules, e.g., round-robin, and synchronous.
 - Set scheduler: enables users to compose custom update schedules.

Tasks Schedulers (2/2)

```
 \begin{array}{l} \textbf{Input: Data Graph } G = (V, E, D) \\ \textbf{Input: Initial task set } \mathcal{T} = \{(f, v_1), (g, v_2), \ldots\} \\ \textbf{while } \mathcal{T} \textit{ is not Empty do} \\ \textbf{1} & \left(f, v\right) \leftarrow \texttt{RemoveNext}\left(\mathcal{T}\right) \\ \textbf{2} & \left(\mathcal{T}', \mathcal{S}_v\right) \leftarrow f(v, \mathcal{S}_v) \\ \textbf{3} & \left(\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'\right) \\ \textbf{Output: Modified Data Graph } G = (V, E, D') \end{array}
```

How to add new task in the queue?

Tasks Schedulers (2/2)

```
Input: Data Graph G = (V, E, D)

Input: Initial task set \mathcal{T} = \{(f, v_1), (g, v_2), ...\}

while \mathcal{T} is not Empty do

1 (f, v) \leftarrow \texttt{RemoveNext}(\mathcal{T})

2 (\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)

3 (\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}')

Output: Modified Data Graph G = (V, E, D')
```

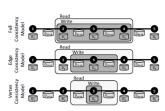
- ► How to add new task in the queue?
 - FIFO: only permits task creation but do not permit task reordering.
 - Prioritized: permits task reordering at the cost of increased overhead.

Consistency

- ▶ Implemented in C++ using PThreads for parallelism.
- ► Consistency: read-write lock

Consistency

- ▶ Implemented in C++ using PThreads for parallelism.
- ► Consistency: read-write lock
- Vertex consistency
 - Central vertex (write-lock)
- Edge consistency
 - Central vertex (write-lock)
 - Adjacent vertices (read-locks)
- ► Full consistency
 - Central vertex (write-locks)
 - Adjacent vertices (write-locks)
- Deadlocks are avoided by acquiring locks sequentially following a canonical order.



GraphLab Implementation

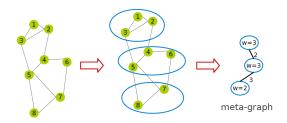
- ► Shared memory implementation
- ► Distributed implementation

Distributed Implementation

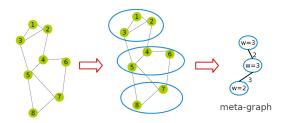
- Graph partitioning
 - How to efficiently load, partition and distribute the data graph across machines?
- ► Consistency
 - How to achieve consistency in the distributed setting?
- ► Fault tolerance

Graph Partitioning

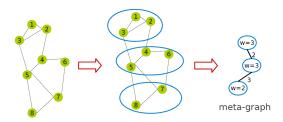
- ► Two-phase partitioning.
- ▶ Partitioning the data graph into k parts, called atom.
 - k ≫ number of machines



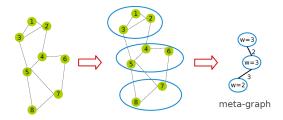
- ► Two-phase partitioning.
- ▶ Partitioning the data graph into k parts, called atom.
 - k ≫ number of machines
- meta-graph: the graph of atoms (one vertex for each atom).

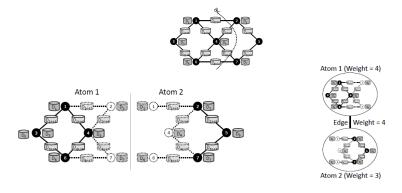


- ► Two-phase partitioning.
- ▶ Partitioning the data graph into k parts, called atom.
 - k ≫ number of machines
- ▶ meta-graph: the graph of atoms (one vertex for each atom).
- ► Atom weight: the amount of data it stores.

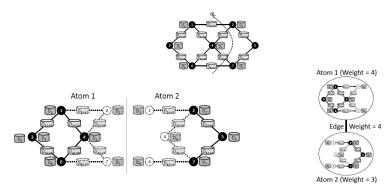


- ► Two-phase partitioning.
- ▶ Partitioning the data graph into k parts, called atom.
 - k ≫ number of machines
- ▶ meta-graph: the graph of atoms (one vertex for each atom).
- ▶ Atom weight: the amount of data it stores.
- ► Edge weight: the number of edges crossing the atoms.

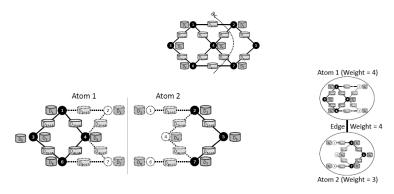




► Each atom is stored as a separate file on a distributed storage system, e.g., HDFS.



- ► Each atom is stored as a separate file on a distributed storage system, e.g., HDFS.
- ► Each atom file is a simple binary that stores interior and the ghosts of the partition information.



- ► Each atom is stored as a separate file on a distributed storage system, e.g., HDFS.
- ► Each atom file is a simple binary that stores interior and the ghosts of the partition information.
- ► Ghost: set of vertices and edges adjacent to the partition boundary.

Graph Partitioning - Phase 2

- Meta-graph is very small.
- ► A fast balanced partition of the meta-graph over the physical machines.
- Assigning graph atoms to machines.

Consistency

Consistency

- ► To achieve a serializable parallel execution of a set of dependent tasks.
 - Chromatic engine
 - Distributed locking engine

► Construct a vertex coloring: assigns a color to each vertex such that no adjacent vertices share the same color.

- Construct a vertex coloring: assigns a color to each vertex such that no adjacent vertices share the same color.
- Edge consistency: executing, synchronously, all update tasks associated with vertices of the same color before proceeding to the next color.

- Construct a vertex coloring: assigns a color to each vertex such that no adjacent vertices share the same color.
- Edge consistency: executing, synchronously, all update tasks associated with vertices of the same color before proceeding to the next color.
- ► Full consistency: no vertex shares the same color as any of its distance two neighbors.

- Construct a vertex coloring: assigns a color to each vertex such that no adjacent vertices share the same color.
- Edge consistency: executing, synchronously, all update tasks associated with vertices of the same color before proceeding to the next color.
- ► Full consistency: no vertex shares the same color as any of its distance two neighbors.
- ► Vertex consistency: assigning all vertices the same color.

► Associating a readers-writer lock with each vertex.

- Associating a readers-writer lock with each vertex.
- Vertex consistency
 - Central vertex (write-lock)

- Associating a readers-writer lock with each vertex.
- Vertex consistency
 - Central vertex (write-lock)
- ► Edge consistency
 - Central vertex (write-lock), Adjacent vertices (read-locks)

- ► Associating a readers-writer lock with each vertex.
- Vertex consistency
 - Central vertex (write-lock)
- ► Edge consistency
 - Central vertex (write-lock), Adjacent vertices (read-locks)
- ► Full consistency
 - Central vertex (write-locks), Adjacent vertices (write-locks)

- ► Associating a readers-writer lock with each vertex.
- Vertex consistency
 - Central vertex (write-lock)
- ► Edge consistency
 - Central vertex (write-lock), Adjacent vertices (read-locks)
- ► Full consistency
 - Central vertex (write-locks), Adjacent vertices (write-locks)
- ► Deadlocks are avoided by acquiring locks sequentially following a canonical order.

Fault Tolerance

► The systems periodically signals all computation activity to halt.

- ► The systems periodically signals all computation activity to halt.
- ► Then synchronizes all caches (ghosts) and saves to disk all data which has been modified since the last snapshot.

- ► The systems periodically signals all computation activity to halt.
- ► Then synchronizes all caches (ghosts) and saves to disk all data which has been modified since the last snapshot.
- Simple, but eliminates the systems advantage of asynchronous computation.

► Based on the Chandy-Lamport algorithm.

- ► Based on the Chandy-Lamport algorithm.
- ► The snapshot function is implemented as an update function in vertices.

- ► Based on the Chandy-Lamport algorithm.
- ► The snapshot function is implemented as an update function in vertices.
- ► The snapshot update takes priority over all other update functions.

- ► Based on the Chandy-Lamport algorithm.
- ► The snapshot function is implemented as an update function in vertices.
- ► The snapshot update takes priority over all other update functions.
- ► Edge Consistency is used on all update functions.

- ► Based on the Chandy-Lamport algorithm.
- ► The snapshot function is implemented as an update function in vertices.
- ► The snapshot update takes priority over all other update functions.
- ► Edge Consistency is used on all update functions.

Mark v as snapshotted

Summary

GraphLab Summary

- ► Asynchronous model
- Vertex-centric
- Communication: distributed shared memory
- ► Three consistency levels: full, edge-level, and vertex-level
- Partitioning: two-phase partitioning
- ► Consistency: chromatic engine (graph coloring), distributed lock engine (reader-writer lock)

GraphLab Limitations

▶ Poor performance on Natural graphs.

Questions?