

PowerGraph and GraphX

Amir H. Payberah
amir@sics.se

Amirkabir University of Technology
(Tehran Polytechnic)



Reminder


Data-Parallel vs. Graph-Parallel Computation

Data-Parallel




hadoop

Spark



Graph-Parallel

Pregel GraphLab



- ▶ Vertex-centric


- ▶ Vertex-centric
- ▶ Bulk Synchronous Parallel (BSP)

- ▶ Vertex-centric
- ▶ Bulk Synchronous Parallel (BSP)
- ▶ Runs in sequence of iterations (supersteps)

- ▶ Vertex-centric
- ▶ Bulk Synchronous Parallel (BSP)
- ▶ Runs in sequence of iterations (supersteps)
- ▶ A vertex in superstep S can:
 - reads messages sent to it in superstep $S-1$.
 - sends messages to other vertices: receiving at superstep $S+1$.
 - modifies its state.

Pregel Limitations

- ▶ **Inefficient** if different regions of the graph converge at **different speed**.
- ▶ Can suffer if one **task** is **more expensive** than the others.
- ▶ Runtime of each phase is determined by the **slowest** machine.



Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

- 1 $(f, v) \leftarrow \text{RemoveNext}(\mathcal{T})$
- 2 $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$
- 3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$



Output: Modified Data Graph $G = (V, E, D')$

GraphLab Limitations

- ▶ Poor performance on **Natural** graphs.


Natural Graphs

- ▶ Graphs derived from **natural phenomena**.
- ▶ Skewed **Power-Law** degree distribution.



Natural Graphs Challenges


- ▶ Traditional graph-partitioning algorithms (**edge-cut** algorithms) perform **poorly** on Power-Law Graphs.
- ▶ Challenges of **high-degree** vertices.




Vertex-Cut Partitioning

Proposed Solution

Vertex-Cut Partitioning




Edge-cut




Vertex-cut


Edge-cut vs. Vertex-cut Partitioning



Edge-cut vs. Vertex-cut Partitioning



Edge-cut



Vertex-cut


PowerGraph

- ▶ Vertex-cut partitioning of graphs.
- ▶ Factorizes the GraphLab update function into the Gather, Apply and Scatter phases (GAS).

Gather-Apply-Scatter Programming Model

► Gather


- **Accumulate** information about neighborhood through a generalized **sum**.



Gather-Apply-Scatter Programming Model

► Gather

- Accumulate information about neighborhood through a generalized sum.




► Apply

- Apply the accumulated value to center vertex.

Gather-Apply-Scatter Programming Model

► Gather

- **Accumulate** information about neighborhood through a generalized **sum**.




► Apply

- **Apply** the accumulated value to center vertex.

► Scatter

- **Update** adjacent edges and vertices.



- ▶ A **directed graph** that stores the program **state**, called **data graph**.

Execution Model (1/2)

- ▶ **Vertex-centric** programming: implementing the **GASVertexProgram** interface (**vertex-program** for short).
- ▶ Similar to **Comput** in **Pregel**, and **update** function in **GraphLab**.

```
interface GASVertexProgram(u) {  
  // Run on gather_nbrs(u)  
  gather( $D_u$ ,  $D_{u-v}$ ,  $D_v$ )  $\rightarrow$  Accum  
  sum(Accum left, Accum right)  $\rightarrow$  Accum  
  apply( $D_u$ , Accum)  $\rightarrow D_u^{\text{new}}$   
  // Run on scatter_nbrs(u)  
  scatter( $D_u^{\text{new}}$ ,  $D_{u-v}$ ,  $D_v$ )  $\rightarrow$  ( $D_{u-v}^{\text{new}}$ , Accum)  
}
```


Execution Model (2/2)

Input: Center vertex u

if *Cache Disabled* **then**

 // Basic Gather-Apply-Scatter Model

foreach neighbor v in $\text{gather_nbrs}(u)$ **do**

$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$

$D_u \leftarrow \text{apply}(D_u, a_u)$

foreach neighbor v in $\text{scatter_nbrs}(u)$ **do**

$(D_{u-v}) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$

else if *Cache Enabled* **then**

 // Faster GAS Model with Delta Caching

if *cached accumulator a_u is empty* **then**

foreach neighbor v in $\text{gather_nbrs}(u)$ **do**

$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$

$D_u \leftarrow \text{apply}(D_u, a_u)$

foreach neighbor v in $\text{scatter_nbrs}(u)$ **do**

$(D_{u-v}, \Delta a) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$

if a_v and Δa are not Empty **then** $a_v \leftarrow \text{sum}(a_v, \Delta a)$

else $a_v \leftarrow \text{Empty}$

Execution Model (2/2)

Input: Center vertex u

if *Cache Disabled* **then**

 // Basic Gather-Apply-Scatter Model

foreach neighbor v in $\text{gather_nbrs}(u)$ **do**

$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$

$D_u \leftarrow \text{apply}(D_u, a_u)$

foreach neighbor v in $\text{scatter_nbrs}(u)$ **do**

$(D_{u-v}) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$

~~**else if** *Cache Enabled* **then**~~

~~// Faster GAS Model with Delta Caching~~

~~**if** *cached accumulator* a_u *is empty* **then**~~

~~**foreach** neighbor v in $\text{gather_nbrs}(u)$ **do**~~

~~$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$~~

~~$D_u \leftarrow \text{apply}(D_u, a_u)$~~

~~**foreach** neighbor v in $\text{scatter_nbrs}(u)$ **do**~~

~~$(D_{u-v}, \Delta a) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$~~

~~**if** a_v and Δa are not Empty **then** $a_v \leftarrow \text{sum}(a_v, \Delta a)$~~

~~**else** $a_v \leftarrow \text{Empty}$~~

PageRank - Pregel

```
Pregel_PageRank(i, messages):  
    // receive all the messages  
    total = 0  
    foreach(msg in messages):  
        total = total + msg  
  
    // update the rank of this vertex  
    R[i] = 0.15 + total  
  
    // send new messages to neighbors  
    foreach(j in out_neighbors[i]):  
        sendmsg(R[i] * wij) to vertex j
```

$$R[i] = 0.15 + \sum_{j \in N_{\text{brs}}(i)} w_{ji} R[j]$$

PageRank - GraphLab

```
GraphLab_PageRank(i)
  // compute sum over neighbors
  total = 0
  foreach(j in in_neighbors(i)):
    total = total + R[j] * wji

  // update the PageRank
  R[i] = 0.15 + total

  // trigger neighbors to run again
  foreach(j in out_neighbors(i)):
    signal vertex-program on j
```

$$R[i] = 0.15 + \sum_{j \in N_{\text{brs}}(i)} w_{ji} R[j]$$

PageRank - PowerGraph

```
PowerGraph_PageRank(i):  
  Gather(j -> i):  
    return wji * R[j]  
  
  sum(a, b):  
    return a + b  
  
  // total: Gather and sum  
  Apply(i, total):  
    R[i] = 0.15 + total  
  
  Scatter(i -> j):  
    if R[i] changed then activate(j)
```

$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

Scheduling (1/5)

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

```
1 |  $(f, v) \leftarrow \text{RemoveNext}(\mathcal{T})$   
2 |  $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$   
3 |  $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$ 
```

Output: Modified Data Graph $G = (V, E, D')$

- PowerGraph inherits the **dynamic scheduling** of **GraphLab**.

Scheduling (2/5)

- Initially all vertices are active.

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

```
1 | (f, v) ← RemoveNext( $\mathcal{T}$ )  
2 | ( $\mathcal{T}'$ ,  $\mathcal{S}_v$ ) ←  $f(v, \mathcal{S}_v)$   
3 |  $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$ 
```

Output: Modified Data Graph $G = (V, E, D')$

Scheduling (2/5)

- Initially **all vertices** are **active**.

Input: Data Graph $G = (V, E, D)$

Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$

while \mathcal{T} is not Empty **do**

```
1 | (f, v) ← RemoveNext( $\mathcal{T}$ )
2 | ( $\mathcal{T}'$ ,  $\mathcal{S}_v$ ) ←  $f(v, \mathcal{S}_v)$ 
3 |  $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$ 
```

Output: Modified Data Graph $G = (V, E, D')$

- PowerGraph executes the **vertex-program** on the **active vertices** until none remain.

Scheduling (2/5)

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty **do**

```
1  |   $(f, v) \leftarrow \text{RemoveNext}(\mathcal{T})$   
2  |   $(\mathcal{T}', \mathcal{S}_v) \leftarrow f(v, \mathcal{S}_v)$   
3  |   $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$ 
```

Output: Modified Data Graph $G = (V, E, D')$

- ▶ Initially all vertices are active.
- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.
- ▶ The order of executing activated vertices is up to the PowerGraph execution engine.

Scheduling (2/5)

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty **do**

```
1   $(f, v) \leftarrow \text{RemoveNext}(\mathcal{T})$   
2   $(\mathcal{T}', S_v) \leftarrow f(v, S_v)$   
3   $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$ 
```

Output: Modified Data Graph $G = (V, E, D')$

- ▶ Initially all vertices are active.
- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.
- ▶ The order of executing activated vertices is up to the PowerGraph execution engine.
- ▶ Once a vertex-program completes the scatter phase it becomes inactive until it is reactivated.

Scheduling (2/5)

Input: Data Graph $G = (V, E, D)$
Input: Initial task set $\mathcal{T} = \{(f, v_1), (g, v_2), \dots\}$
while \mathcal{T} is not Empty **do**
1 $(f, v) \leftarrow \text{RemoveNext}(\mathcal{T})$
2 $(\mathcal{T}', S_v) \leftarrow f(v, S_v)$
3 $\mathcal{T} \leftarrow \mathcal{T} \cup \mathcal{T}'$
Output: Modified Data Graph $G = (V, E, D')$

- ▶ Initially all vertices are active.
- ▶ PowerGraph executes the vertex-program on the active vertices until none remain.
- ▶ The order of executing activated vertices is up to the PowerGraph execution engine.
- ▶ Once a vertex-program completes the scatter phase it becomes inactive until it is reactivated.
- ▶ Vertices can activate themselves and neighboring vertices.

- ▶ PowerGraph can execute both **synchronously** and **asynchronously**.
 - Bulk synchronous execution
 - Asynchronous execution

Scheduling - Bulk Synchronous Execution (4/5)

- ▶ Similar to Pregel.

Scheduling - Bulk Synchronous Execution (4/5)

- ▶ Similar to Pregel.
- ▶ Minor-step: executing the gather, apply, and scatter in order.
 - Runs synchronously on all active vertices with a barrier at the end.

Scheduling - Bulk Synchronous Execution (4/5)

- ▶ Similar to Pregel.
- ▶ **Minor-step**: executing the gather, apply, and scatter in order.
 - Runs synchronously on all active vertices with a barrier at the end.
- ▶ **Super-step**: a complete series of GAS minor-steps.

Scheduling - Bulk Synchronous Execution (4/5)

- ▶ Similar to Pregel.
- ▶ **Minor-step**: executing the **gather, apply, and scatter in order**.
 - Runs **synchronously** on all **active** vertices with a **barrier** at the end.
- ▶ **Super-step**: a complete series of GAS minor-steps.
- ▶ Changes made to the vertex/edge data are committed at the **end** of each **minor-step** and are visible in the **subsequent minor-steps**.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply and scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply and scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply and scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.
 - **Dining philosophers** problem, where each vertex is a philosopher, and each edge is a fork.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply and scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.
 - **Dining philosophers** problem, where each vertex is a philosopher, and each edge is a fork.
 - **GraphLab** implements **Dijkstras** solution, where forks are acquired **sequentially** according to a total ordering.

Scheduling - Asynchronous Execution (5/5)

- ▶ Changes made to the vertex/edge data during the **apply and scatter** functions are **immediately** committed to the graph.
 - **Visible** to subsequent computation on neighboring vertices.
- ▶ **Serializability**: prevents **adjacent vertex-programs** from running **concurrently** using a **fine-grained locking** protocol.
 - **Dining philosophers** problem, where each vertex is a philosopher, and each edge is a fork.
 - **GraphLab** implements **Dijkstras** solution, where forks are acquired **sequentially** according to a total ordering.
 - **PowerGraph** implements **Chandy-Misra** solution, which acquires all forks **simultaneously**.

Delta Caching (1/2)

- ▶ Changes in a few of its neighbors → triggering a vertex-program
- ▶ The gather operation is invoked on all neighbors: wasting computation cycles

Delta Caching (1/2)

- ▶ Changes in a few of its neighbors \rightarrow triggering a vertex-program
- ▶ The gather operation is invoked on all neighbors: wasting computation cycles
- ▶ Maintaining a cache of the accumulator a_v from the previous gather phase for each vertex.
- ▶ The scatter can return an additional Δa , which is added to the cached accumulator a_v .

Delta Caching (2/2)

Input: Center vertex u

if *Cache Disabled* **then**

 // Basic Gather-Apply-Scatter Model

foreach neighbor v in $\text{gather_nbrs}(u)$ **do**

$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$

$D_u \leftarrow \text{apply}(D_u, a_u)$

foreach neighbor v in $\text{scatter_nbrs}(u)$ **do**

$(D_{u-v}) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$

~~**else if** *Cache Enabled* **then**~~

~~// Faster GAS Model with Delta Caching~~

~~**if** *cached accumulator* a_u *is empty* **then**~~

~~**foreach** neighbor v in $\text{gather_nbrs}(u)$ **do**~~

~~$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$~~

~~$D_u \leftarrow \text{apply}(D_u, a_u)$~~

~~**foreach** neighbor v in $\text{scatter_nbrs}(u)$ **do**~~

~~$(D_{u-v}, \Delta a) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$~~

~~**if** a_v and Δa are not Empty **then** $a_v \leftarrow \text{sum}(a_v, \Delta a)$~~

~~**else** $a_v \leftarrow \text{Empty}$~~

Delta Caching (2/2)

Input: Center vertex u

if *Cache Disabled* **then**

 // Basic Gather-Apply-Scatter Model

foreach neighbor v in $\text{gather_nbrs}(u)$ **do**

$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$

$D_u \leftarrow \text{apply}(D_u, a_u)$

foreach neighbor v scatter_nbrs(u) **do**

$(D_{u-v}) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$

else if *Cache Enabled* **then**

 // Faster GAS Model with Delta Caching

if cached accumulator a_u is empty **then**

foreach neighbor v in $\text{gather_nbrs}(u)$ **do**

$a_u \leftarrow \text{sum}(a_u, \text{gather}(D_u, D_{u-v}, D_v))$

$D_u \leftarrow \text{apply}(D_u, a_u)$

foreach neighbor v scatter_nbrs(u) **do**

$(D_{u-v}, \Delta a) \leftarrow \text{scatter}(D_u, D_{u-v}, D_v)$

if a_v and Δa are not Empty **then** $a_v \leftarrow \text{sum}(a_v, \Delta a)$

else $a_v \leftarrow \text{Empty}$

Example: PageRank (Delta-Caching)

```
PowerGraph_PageRank(i):  
  Gather(j -> i):  
    return wji * R[j]  
  
sum(a, b):  
  return a + b  
  
// total: Gather and sum  
Apply(i, total):  
  new = 0.15 + total  
  R[i].delta = new - R[i]  
  R[i] = new  
  
Scatter(i -> j):  
  if R[i] changed then activate(j)  
  return wij * R[i].delta
```

$$R[i] = 0.15 + \sum_{j \in Nbrs(i)} w_{ji} R[j]$$

Graph Partitioning

- ▶ Vertex-cut partitioning.
- ▶ Evenly assign edges to machines.
 - Minimize machines spanned by each vertex.
- ▶ Two proposed solutions:
 - Random edge placement.
 - Greedy edge placement.

Random Vertex-Cuts

- ▶ Randomly assign edges to machines.
- ▶ Completely parallel and easy to distribute.
- ▶ High replication factor.

Greedy Vertex-Cuts (1/2)

- ▶ $A(v)$: set of machines that contain adjacent edges of v .

Greedy Vertex-Cuts (1/2)

- ▶ $A(v)$: set of machines that contain adjacent edges of v .
- ▶ **Case 1**: If $A(u)$ and $A(v)$ intersect, then the edge should be assigned to a machine in the intersection.

Greedy Vertex-Cuts (1/2)

- ▶ $A(v)$: set of machines that contain adjacent edges of v .
- ▶ **Case 1**: If $A(u)$ and $A(v)$ intersect, then the edge should be assigned to a machine in the intersection.
- ▶ **Case 2**: If $A(u)$ and $A(v)$ are not empty and do not intersect, then the edge should be assigned to one of the machines from the vertex with the most unassigned edges.

Greedy Vertex-Cuts (1/2)

- ▶ $A(v)$: set of machines that contain adjacent edges of v .
- ▶ **Case 1:** If $A(u)$ and $A(v)$ intersect, then the edge should be assigned to a machine in the intersection.
- ▶ **Case 2:** If $A(u)$ and $A(v)$ are not empty and do not intersect, then the edge should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ **Case 3:** If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.

Greedy Vertex-Cuts (1/2)

- ▶ $A(v)$: set of machines that contain adjacent edges of v .
- ▶ **Case 1**: If $A(u)$ and $A(v)$ intersect, then the edge should be assigned to a machine in the intersection.
- ▶ **Case 2**: If $A(u)$ and $A(v)$ are not empty and do not intersect, then the edge should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ **Case 3**: If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.
- ▶ **Case 4**: If neither vertex has been assigned, then assign the edge to the least loaded machine.

Greedy Vertex-Cuts (2/2)

- ▶ **Coordinated** edge placement:
 - Requires coordination to place each edge
 - **Slower**, but **higher** quality cuts
- ▶ **Oblivious** edge placement:
 - Approx. greedy objective without coordination
 - **Faster**, but **lower** quality cuts

PowerGraph Summary

- ▶ Gather-Apply-Scatter programming model
- ▶ Synchronous and Asynchronous models
- ▶ Vertex-cut graph partitioning

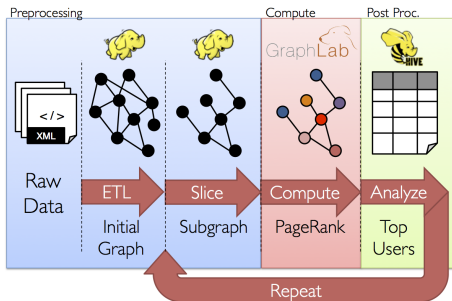
► Any limitations?

Data-Parallel vs. Graph-Parallel Computation


- ▶ **Graph-parallel** computation: **restricting** the types of computation to achieve **performance**.

Data-Parallel vs. Graph-Parallel Computation

- ▶ **Graph-parallel** computation: **restricting** the types of computation to achieve **performance**.
- ▶ **But**, the same restrictions make it **difficult** and **inefficient** to express many stages in a typical graph-analytics **pipeline**.




Data-Parallel and Graph-Parallel Pipeline

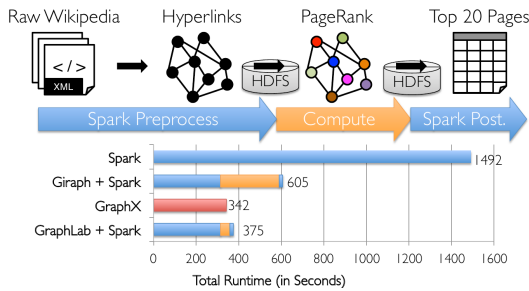




- ▶ **Moving** between **table** and **graph** views of the **same physical data**.
- ▶ **Inefficient**: extensive **data movement** and **duplication** across the network and file system.

GraphX vs. Data-Parallel/Graph-Parallel Systems



GraphX vs. Data-Parallel/Graph-Parallel Systems






- ▶ New API that blurs the distinction between Tables and Graphs.
- ▶ New system that unifies Data-Parallel and Graph-Parallel systems.
- ▶ It is implemented on top of Spark.

Unifying Data-Parallel and Graph-Parallel Analytics

- ▶ **Tables** and **Graphs** are **composable views** of the same physical data.
- ▶ Each view has its **own operators** that **exploit the semantics** of the view to achieve **efficient** execution.




Data Model

► **Property Graph**: represented using **two** Spark **RDDs**:

- **Edge collection**: VertexRDD
- **Vertex collection**: EdgeRDD

```
// VD: the type of the vertex attribute
// ED: the type of the edge attribute
class Graph[VD, ED] {
  val vertices: VertexRDD[VD]
  val edges: EdgeRDD[ED]
}
```




Primitive Data Types

```
// Vertex collection
class VertexRDD[VD] extends RDD[(VertexId, VD)]


// Edge collection
class EdgeRDD[ED] extends RDD[Edge[ED]]
case class Edge[ED](srcId: VertexId = 0, dstId: VertexId = 0,
                    attr: ED = null.asInstanceOf[ED])

// Edge Triple
class EdgeTriplet[VD, ED] extends Edge[ED]
```

- ▶ **EdgeTriplet** represents an **edge** along with the **vertex attributes** of its **neighboring** vertices.



Example (1/3)



Vertex Table

Id	Property (V)
3	(rxin, student)
7	(jgonzal, postdoc)
5	(franklin, professor)
2	(istoica, professor)

Edge Table

SrcId	DstId	Property (E)
3	7	Collaborator
5	3	Advisor
2	5	Colleague
5	7	PI

Example (2/3)

```
val sc: SparkContext

// Create an RDD for the vertices
val users: VertexRDD[(String, String)] = sc.parallelize(
    Array((3L, ("rxin", "student")), (7L, ("jgonzal", "postdoc")),
          (5L, ("franklin", "prof")), (2L, ("istoica", "prof"))))

// Create an RDD for edges
val relationships: EdgeRDD[String] = sc.parallelize(
    Array(Edge(3L, 7L, "collab"),    Edge(5L, 3L, "advisor"),
          Edge(2L, 5L, "colleague"), Edge(5L, 7L, "pi")))

// Define a default user in case there are relationship with missing user
val defaultUser = ("John Doe", "Missing")

// Build the initial Graph
val userGraph: Graph[(String, String), String] =
    Graph(users, relationships, defaultUser)
```


Example (3/3)

```
// Constructed from above
val userGraph: Graph[(String, String), String]

// Count all users which are postdocs
userGraph.vertices.filter((id, (name, pos)) => pos == "postdoc").count

// Count all the edges where src > dst
userGraph.edges.filter(e => e.srcId > e.dstId).count

// Use the triplets view to create an RDD of facts
val facts: RDD[String] = graph.triplets.map(triplet =>
    triplet.srcAttr._1 + " is the " +
    triplet.attr + " of " + triplet.dstAttr._1)

// Remove missing vertices as well as the edges to connected to them
val validGraph = graph.subgraph(vpred = (id, attr) => attr._2 != "Missing")

facts.collect.foreach(println(_))
```

Property Operators (1/2)

```
class Graph[VD, ED] {  
  def mapVertices[VD2](map: (VertexId, VD) => VD2): Graph[VD2, ED]  
  
  def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]  
  
  def mapTriplets[ED2](map: EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]  
}
```

- ▶ They yield **new graphs** with the vertex or edge properties modified by the map function.
- ▶ The graph **structure** is **unaffected**.

Property Operators (2/2)



```
val newGraph = graph.mapVertices((id, attr) => mapUdf(id, attr))
```

```
val newVertices = graph.vertices.map((id, attr) => (id, mapUdf(id, attr)))  
val newGraph = Graph(newVertices, graph.edges)
```

- ▶ Both are logically equivalent, but the second one **does not preserve** the structural indices and would not benefit from the GraphX system **optimizations**.



Map Reduce Triplets

- ▶ Map-Reduce for each vertex



Map Reduce Triplets

- ▶ Map-Reduce for each vertex



```
// what is the age of the oldest follower for each user?  
val oldestFollowerAge = graph.mapReduceTriplets(  
  e => (e.dstAttr, e.srcAttr), // Map  
  (a, b) => max(a, b) // Reduce  
)
```

Structural Operators

```
class Graph[VD, ED] {  
  // returns a new graph with all the edge directions reversed  
  def reverse: Graph[VD, ED]  
  
  // returns the graph containing only the vertices and edges that satisfy  
  // the vertex predicate  
  def subgraph(epred: EdgeTriplet[VD,ED] => Boolean,  
               vpred: (VertexId, VD) => Boolean): Graph[VD, ED]  
  
  // a subgraph by returning a graph that contains the vertices and edges  
  // that are also found in the input graph  
  def mask[VD2, ED2](other: Graph[VD2, ED2]): Graph[VD, ED]  
}
```

Structural Operators Example

```
// Build the initial Graph  
val graph = Graph(users, relationships, defaultUser)  
  
// Run Connected Components  
val ccGraph = graph.connectedComponents()  
  
// Remove missing vertices as well as the edges to connected to them  
val validGraph = graph.subgraph(vpred = (id, attr) => attr._2 != "Missing")  
  
// Restrict the answer to the valid subgraph  
val validCCGraph = ccGraph.mask(validGraph)
```

Join Operators

- To join data from external collections (RDDs) with graphs.

```
class Graph[VD, ED] {  
  // joins the vertices with the input RDD and returns a new graph  
  // by applying the map function to the result of the joined vertices  
  def joinVertices[U](table: RDD[(VertexId, U)])  
    (map: (VertexId, VD, U) => VD): Graph[VD, ED]  
  
  // similarly to joinVertices, but the map function is applied to  
  // all vertices and can change the vertex property type  
  def outerJoinVertices[U, VD2](table: RDD[(VertexId, U)])  
    (map: (VertexId, VD, Option[U]) => VD2): Graph[VD2, ED]  
}
```


Graph Builders

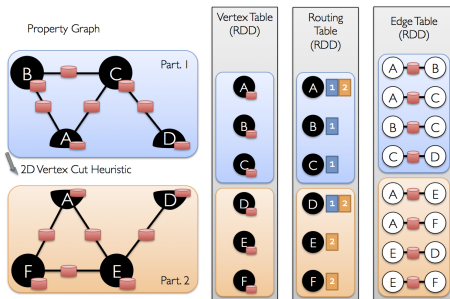
```
// load a graph from a list of edges on disk
object GraphLoader {
  def edgeListFile(
    sc: SparkContext,
    path: String,
    canonicalOrientation: Boolean = false,
    minEdgePartitions: Int = 1)
    : Graph[Int, Int]
}

// graph file
# This is a comment
2 1
4 1
1 2
```

- ▶ GraphX is implemented on top of Spark
- ▶ In-memory caching
- ▶ Lineage-based fault tolerance
- ▶ Programmable partitioning

Distributed Graph Representation (1/2)

- ▶ Representing graphs using **two RDDs**: **edge-collection** and **vertex-collection**
- ▶ **Vertex-cut** partitioning (like **PowerGraph**)



Distributed Graph Representation (2/2)

- ▶ Each vertex partition contains a **bitmask** and **routing table**.
- ▶ **Routing table**: a **logical map** from a vertex id to the set of edge partitions that contains adjacent edges.
- ▶ **Bitmask**: enables the set intersection and filtering.
 - Vertices bitmasks are updated after each operation (e.g., mapReduceTriplets).
 - Vertices hidden by the bitmask **do not** participate in the graph operations.

Summary

► PowerGraph

- GAS programming model
- Vertex-cut partitioning

► GraphX

- Unifying data-parallel and graph-parallel analytics
- Vertex-cut partitioning

Questions?

Acknowledgements

Some pictures were derived from the Spark web site
(<http://spark.apache.org/>).

Questions?