

# Introduction to Operating Systems (Part I)

Amir H. Payberah  
[amir@sics.se](mailto:amir@sics.se)

Amirkabir University of Technology  
(Tehran Polytechnic)



# Course Information

## Course Objective

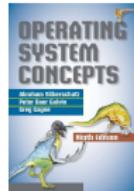
- ▶ The purpose of this course is to teach the design of operating systems.

# Course Objective

- ▶ The purpose of this course is to teach the design of **operating systems**.
- ▶ Topics we will cover include:
  - **Process** management
  - **Memory** management
  - **File systems**
  - **I/O** management
  - **Security** and **privacy**

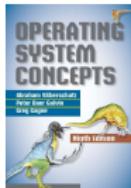
# Course Textbooks

- ▶ Operating System Concepts, 9th Edition  
Avil Silberschatz et al., Wiley, 2013

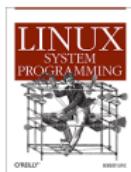


# Course Textbooks

- ▶ Operating System Concepts, 9th Edition  
Avil Silberschatz et al., Wiley, 2013



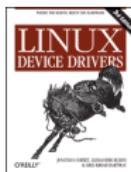
- ▶ Linux System Programming, 2nd Edition  
Robert Love, O'Reilly Media, 2013



- ▶ The Linux Programming Interface  
Michael Kerrisk, No Starch Press, 2010



- ▶ Linux Device Drivers, 3rd Edition  
Jonathan Corbet et al., O'Reilly Media, 2005



# Course Examination

- ▶ Mid term exam: 30%
- ▶ Final exam: 30%
- ▶ Lab assignments: 40%
  - Seven programming assignments in C
  - Students will work in groups of three

- ▶ You can find all the course information on the course web page:  
<http://www.sics.se/~amir/os14.htm>
- ▶ Use the course **discussion forum** if you have any questions.

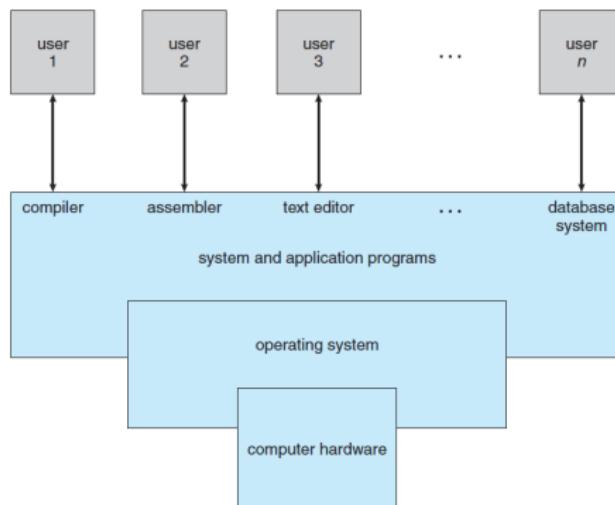
# What is an Operating System?

# What is an Operating System?

- ▶ A **program** that acts as an **intermediary** between a **user** of a computer and the computer **hardware**.

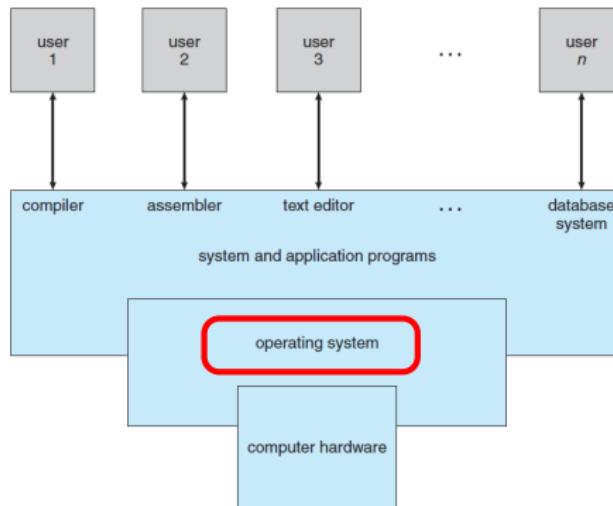
# What is an Operating System?

- ▶ A **program** that acts as an **intermediary** between a **user** of a computer and the computer **hardware**.



# What is an Operating System?

- ▶ A **program** that acts as an **intermediary** between a **user** of a computer and the computer **hardware**.



# Operating System Goals

- ▶ Execute user programs and make solving user problems easier.

# Operating System Goals

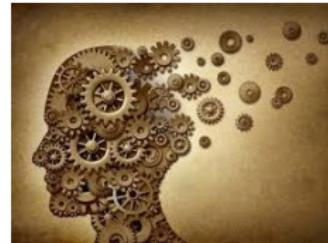
- ▶ Execute user programs and make solving user problems easier.
- ▶ Make the computer system convenient to use.

# Operating System Goals

- ▶ Execute user programs and make solving user problems easier.
- ▶ Make the computer system convenient to use.
- ▶ Use the computer hardware in an efficient manner.

# What Operating Systems Do

- ▶ OS is a **resource allocator**
  - **Manages** all resources.
  - Decides between **conflicting requests** for **efficient** and **fair** resource use.



# What Operating Systems Do

- ▶ OS is a **resource allocator**
  - Manages all resources.
  - Decides between conflicting requests for efficient and fair resource use.
- ▶ OS is a **program controller**
  - Controls execution of programs to prevent errors and improper use of the computer.



# Operating Systems Definition

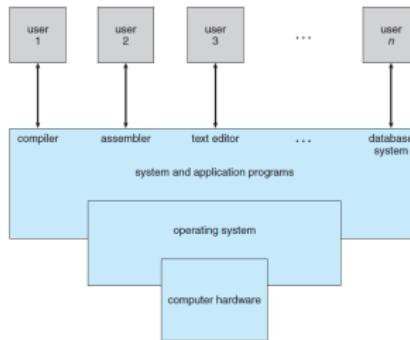
- ▶ No universally accepted definition.

# Operating Systems Definition

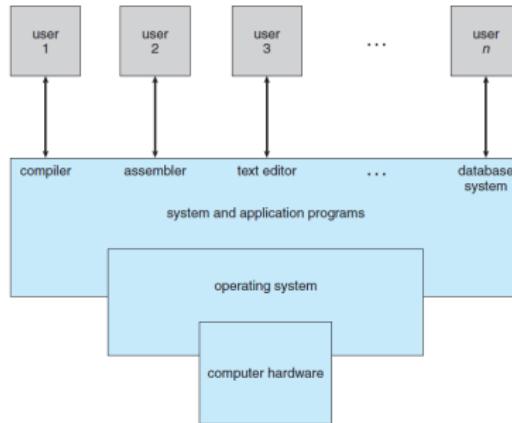
- ▶ No universally accepted definition.
- ▶ The operating system is the one program running at all times on the computer, usually called the kernel.

# Operating Systems Definition

- ▶ No universally accepted definition.
- ▶ The operating system is the one program running at all times on the computer, usually called the **kernel**.
- ▶ Everything else is either a **system program** or an **application program**.



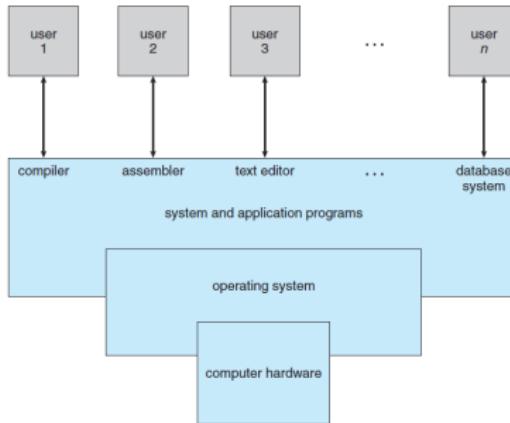
# Four Components of a Computer System (1/4)



## ► Hardware

- Provides basic computing resources.
- CPU, memory, I/O devices, ...

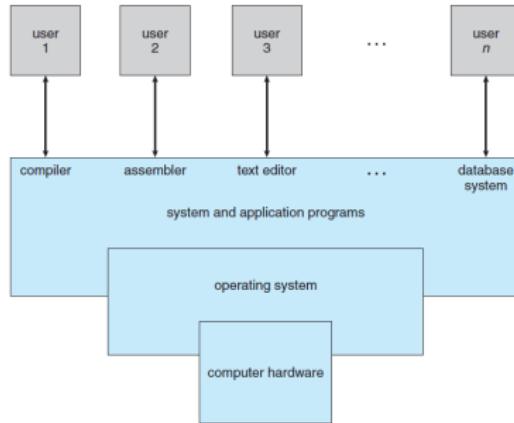
## Four Components of a Computer System (2/4)



### ► Operating system

- Controls and coordinates use of hardware among various applications and users.

# Four Components of a Computer System (3/4)



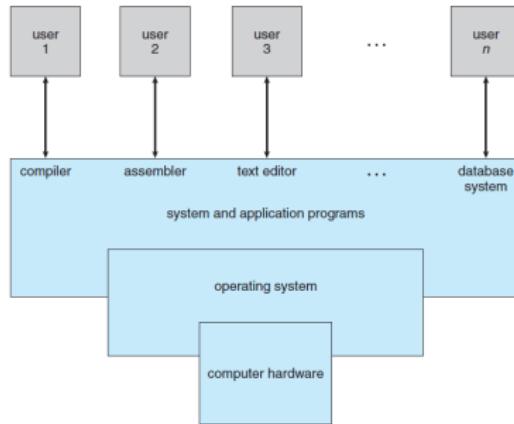
## ► Application programs

- Compilers, web browsers, database systems, video games, ...

## ► System programs

- File manipulation, program loading and execution, ...

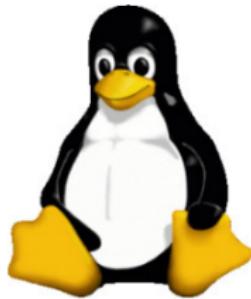
# Four Components of a Computer System (4/4)



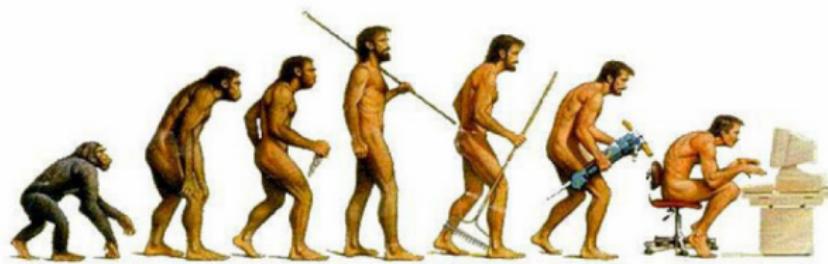
## ► Users

- People, machines, other computers

# Operating Systems

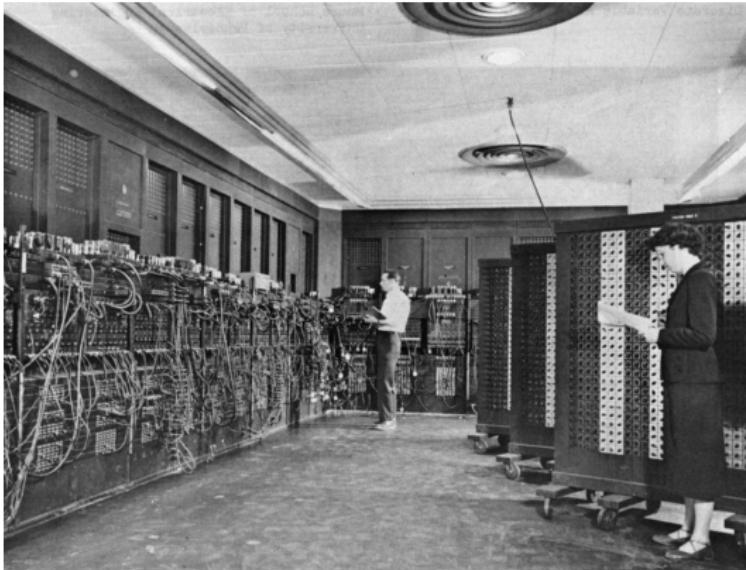


# A Brief History of Operating Systems



# First Generation: 1945-1955 (1/3)

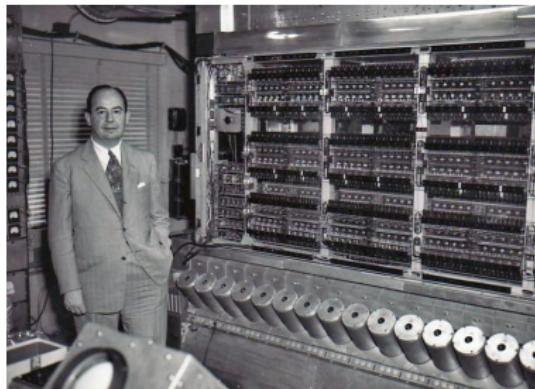
- ▶ No operating system
- ▶ Vacuum tubes and plugboards



ENIAC (Electronic Numerical Integrator And Computer): the first electronic general-purpose computer.  
[\[http://en.wikipedia.org/wiki/ENIAC\]](http://en.wikipedia.org/wiki/ENIAC)

## First Generation: 1945-1955 (2/3)

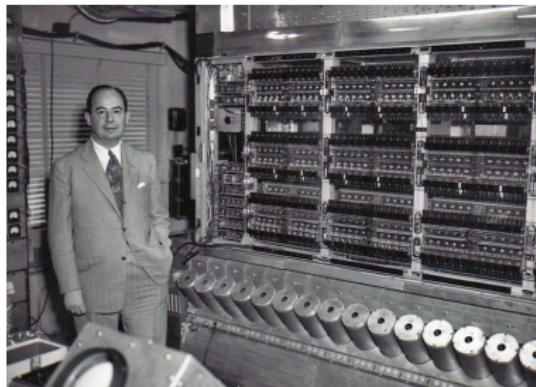
- ▶ Human was the operator and programmer.



John von Neumann  
[<http://ysfine.com/wigner/neumann.html>]

## First Generation: 1945-1955 (2/3)

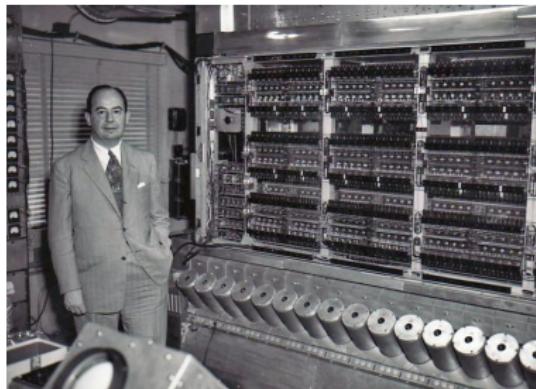
- ▶ Human was the operator and programmer.
- ▶ Computer were programmed by physically re-wiring it; later, through stored programs (von Neumann architecture).



John von Neumann  
[<http://ysfine.com/wigner/neumann.html>]

## First Generation: 1945-1955 (2/3)

- ▶ Human was the operator and programmer.
- ▶ Computer were programmed by physically re-wiring it; later, through stored programs (von Neumann architecture).
- ▶ Programs written in machine or assembly language.



John von Neumann  
[\[http://ysfine.com/wigner/neumann.html\]](http://ysfine.com/wigner/neumann.html)

## First Generation: 1945-1955 (3/3)

- ▶ Problems:

- **Serial processing**: users had access to the computer **one by one** in series.
- Users have to write **again and again** the same routines.

## Second Generation: 1955-1965 (1/5)

- ▶ Transistors
- ▶ Mainframes



IBM 7094 at Columbia University  
[<http://www.columbia.edu/cu/computinghistory/1965.html>]

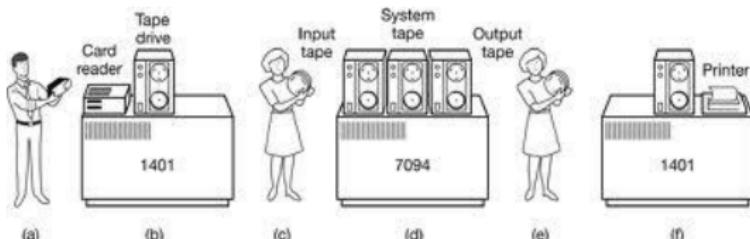
## Second Generation: 1955-1965 (2/5)

- ▶ Separation between **operators** and **programmers**.
  - The **programmer**: prepares her/his **job** off-line.
  - The **operator**: runs the job and delivers a printed output.

- ▶ Separation between **operators** and **programmers**.
  - The **programmer**: prepares her/his **job** off-line.
  - The **operator**: runs the job and delivers a printed output.
- ▶ **Job**
  - A program or set of programs
  - A programmer would first write the program on paper (in FORTRAN or assembly), then **punch it on cards**.

## Second Generation: 1955-1965 (3/5)

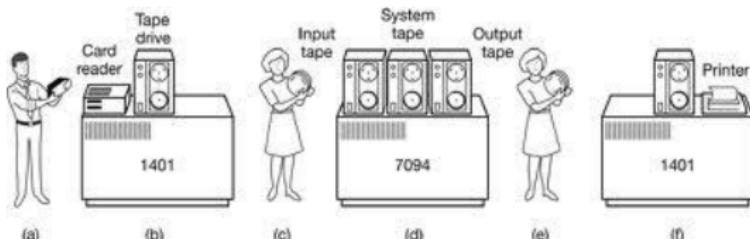
- ▶ Batch the jobs together.



[A.S. Tanenbaum et al., Operating Systems Design and Implementation, 2006]

## Second Generation: 1955-1965 (3/5)

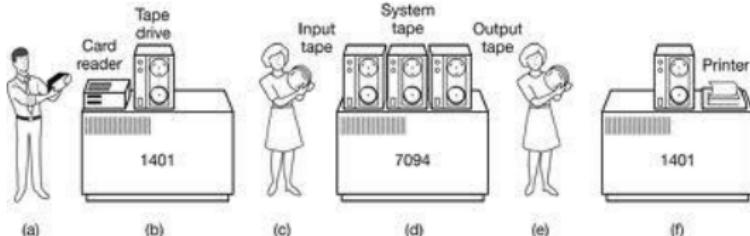
- ▶ Batch the jobs together.
- ▶ The operator pre-reads jobs onto a magnetic tape.



[A.S. Tanenbaum et al., Operating Systems Design and Implementation, 2006]

## Second Generation: 1955-1965 (3/5)

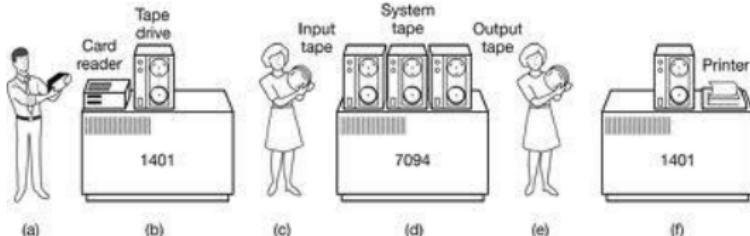
- ▶ Batch the jobs together.
- ▶ The operator pre-reads jobs onto a magnetic tape.
- ▶ The operator loads a special program (monitor) that reads the jobs from the tape and run them sequentially.



[A.S. Tanenbaum et al., Operating Systems Design and Implementation, 2006]

## Second Generation: 1955-1965 (3/5)

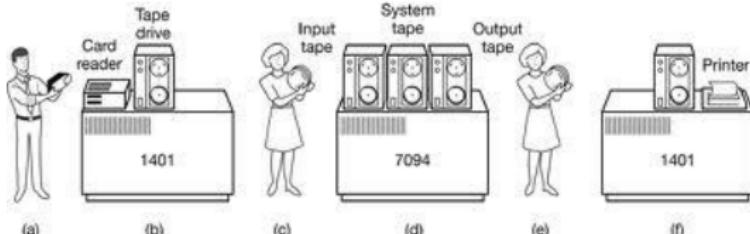
- ▶ Batch the jobs together.
- ▶ The operator pre-reads jobs onto a magnetic tape.
- ▶ The operator loads a special program (monitor) that reads the jobs from the tape and run them sequentially.
- ▶ The monitor program writes the output of each job on a second magnetic tape.



[A.S. Tanenbaum et al., Operating Systems Design and Implementation, 2006]

## Second Generation: 1955-1965 (3/5)

- ▶ Batch the jobs together.
- ▶ The operator pre-reads jobs onto a magnetic tape.
- ▶ The operator loads a special program (monitor) that reads the jobs from the tape and run them sequentially.
- ▶ The monitor program writes the output of each job on a second magnetic tape.
- ▶ The operator brings the full output tape for offline printing.



[A.S. Tanenbaum et al., Operating Systems Design and Implementation, 2006]

## Second Generation: 1955-1965 (4/5)

### ► Monitor

- It is the ancestor of modern O/S.

### ► Monitor

- It is the ancestor of modern O/S.
- A **special program** that **controls** the **sequence of events**.

### ► Monitor

- It is the ancestor of modern O/S.
- A **special program** that **controls** the **sequence of events**.
- Always **resides in main memory**.

### ► Monitor

- It is the ancestor of modern O/S.
- A **special program** that **controls** the **sequence of events**.
- Always **resides in main memory**.
- **Reads in jobs one at a time**, places a job in the **user program area** of the memory, and passes control to it.

### ► Monitor

- It is the ancestor of modern O/S.
- A **special program** that **controls** the **sequence of events**.
- Always **resides in main memory**.
- **Reads in jobs one at a time**, places a job in the **user program area** of the memory, and passes control to it.
- Upon **completion**, the user program branches back to the monitor, which immediately loads and executes the next job

## Second Generation: 1955-1965 (5/5)

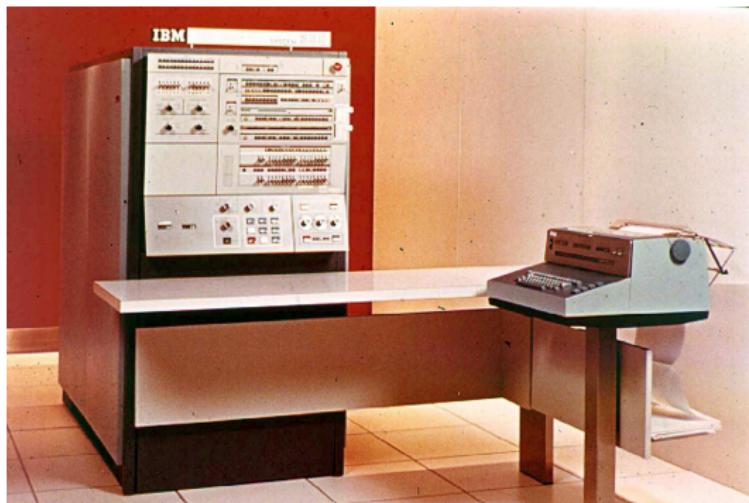
### ► Problems:

- A lot of CPU time is still wasted waiting for I/O instructions to complete.
- I/O devices much slower than processor.



# Third Generation: 1965-1980 (1/5)

- ▶ Integrated Circuits (ICs)

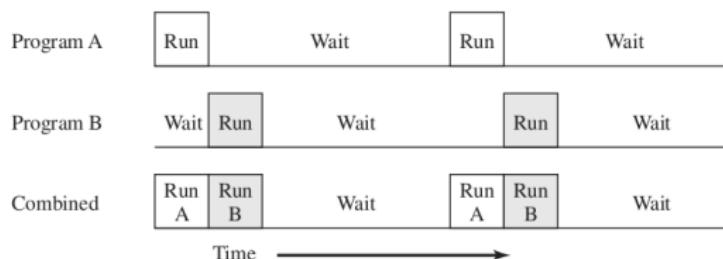


IBM 360

[<http://www.computermuseum.li/Testpage/IBM-360-1964.htm>]

## Third Generation: 1965-1980 (2/5)

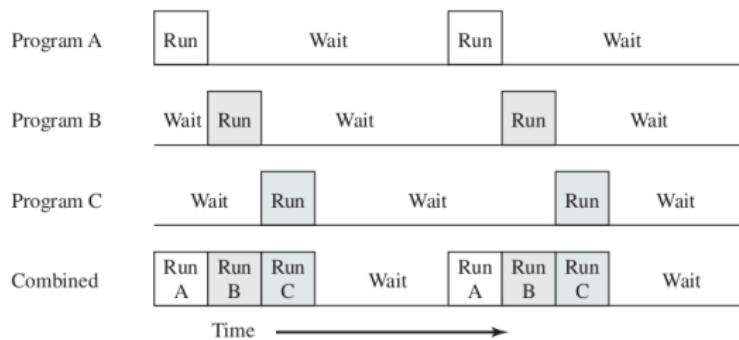
- ▶ Multiprogrammed batch systems.
- ▶ Load two jobs in memory: while one job is waiting for I/O, the processor could switch to the other job.



[W. Stallings, Operating Systems: Internals and Design Principles, 2011]

## Third Generation: 1965-1980 (3/5)

- ▶ Expand to three, four or more jobs.
- ▶ Jobs are kept in **main memory** at the same time and the CPU is multiplexed among them or **multiprogrammed**.



[W. Stallings, Operating Systems: Internals and Design Principles, 2011]

## Third Generation: 1965-1980 (4/5)

- ▶ Multiprogramming alone does not give any guarantee that a program will run in a timely manner.

## Third Generation: 1965-1980 (4/5)

- ▶ Multiprogramming alone does not give any guarantee that a program will run in a timely manner.
- ▶ Tasks kept running until they performed an operation that required waiting for an external event such as I/O.

## Third Generation: 1965-1980 (4/5)

- ▶ Multiprogramming alone does not give any guarantee that a program will run in a timely manner.
- ▶ Tasks kept running until they performed an operation that required waiting for an external event such as I/O.
- ▶ But, in a multiple-user system, users want to see their program running as if it was the only program in the computer.

## Third Generation: 1965-1980 (4/5)

- ▶ Multiprogramming alone does not give any guarantee that a program will run in a timely manner.
- ▶ Tasks kept running until they performed an operation that required waiting for an external event such as I/O.
- ▶ But, in a multiple-user system, users want to see their program running as if it was the only program in the computer.
- ▶ Solution? time-sharing or preemptive multitasking systems.

## ► Time-sharing

- A logical extension of multiprogramming for handling **multiple interactive jobs** among **multiple users**.
- Hardware **timer interrupt**: switching jobs.

## Third Generation: 1965-1980 (5/5)

### ► Time-sharing

- A logical extension of multiprogramming for handling **multiple interactive jobs** among **multiple users**.
- Hardware **timer interrupt**: switching jobs.

### ► Birth of **UNIX**: CTSS → MULTICS → UNIX

**UNIX®**

## Fourth Generation: 1980-Present (1/3)

- ▶ Personal Computers (PCs)



## Fourth Generation: 1980-Present (2/3)

- ▶ From multiple users back to a **single user**.
- ▶ **Multitasking** a central feature of modern PC operating systems.
- ▶ PC systems emphasize **user convenience**.

## Fourth Generation: 1980-Present (3/3)

- ▶ GNU (GNU's Not Unix!): 1983



- ▶ Mac OS: 1984



- ▶ Microsoft Windows: 1985



- ▶ Linux: 1991



## GNU/Linux (1/4)

- ▶ In 1971 Richard Matthew Stallman (RMS) joined MIT.



## GNU/Linux (1/4)

- ▶ In 1971 Richard Matthew Stallman (RMS) joined MIT.
- ▶ At that time, all the programmers used to share their code freely.



## GNU/Linux (1/4)

- ▶ In 1971 Richard Matthew Stallman (RMS) joined MIT.
- ▶ At that time, all the programmers used to share their code freely.
- ▶ In 1980, software companies refused to share the code (copyright).



# GNU/Linux (1/4)

- ▶ In 1971 Richard Matthew Stallman (RMS) joined MIT.
- ▶ At that time, all the programmers used to share their code freely.
- ▶ In 1980, software companies refused to share the code (copyright).
- ▶ In 1985, in response, Stallman, founded the Free Software Foundation (FSF) and published the GNU manifesto.
  - Outlined his motivation for creating a free OS (GNU), which would be compatible with Unix.



## GNU/Linux (2/4)

- ▶ In 1989, Stallman released the first program independent GNU General Public License (GPL) or copyleft.
- ▶ Now the only thing that GNU lacked was a completely free OS kernel: GNU Hurd kernel



## GNU/Linux (3/4)

- ▶ In 1985, [Andy Tanenbaum](#) wrote a [Unix like OS](#) from scratch, called [Minix](#).
- ▶ He implemented it for [educational purposes](#).



## GNU/Linux (4/4)

- ▶ In 1990, [Linus Torvalds](#) wanted to improve Minix.



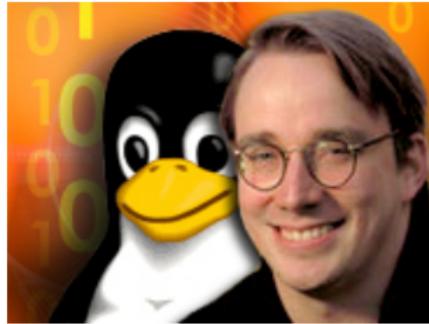
## GNU/Linux (4/4)

- ▶ In 1990, **Linus Torvalds** wanted to improve Minix.
- ▶ But he was **prohibited** by Tanenbaum to do so.



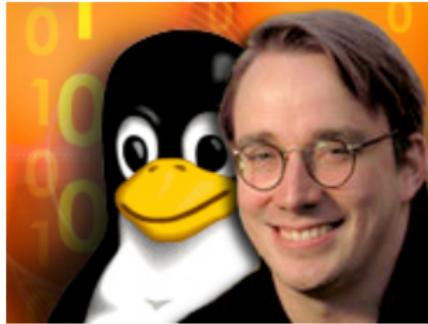
## GNU/Linux (4/4)

- ▶ In 1990, Linus Torvalds wanted to improve Minix.
- ▶ But he was prohibited by Tanenbaum to do so.
- ▶ So, Linus implemented his own kernel and released it under GPL: Linux kernel



## GNU/Linux (4/4)

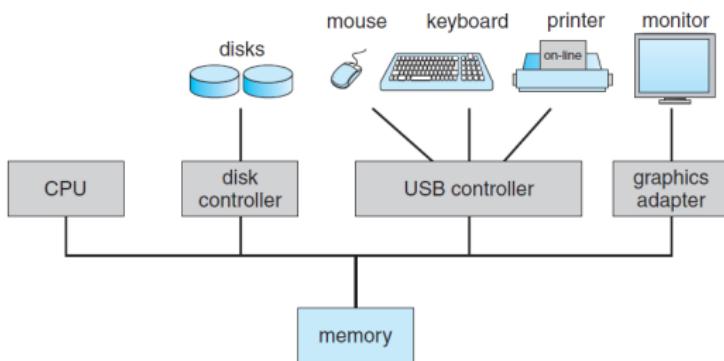
- ▶ In 1990, Linus Torvalds wanted to improve Minix.
- ▶ But he was prohibited by Tanenbaum to do so.
- ▶ So, Linus implemented his own kernel and released it under GPL: Linux kernel
- ▶ Linux, is then, used as the kernel of the GNU in many distributions.



# Computer System Organization

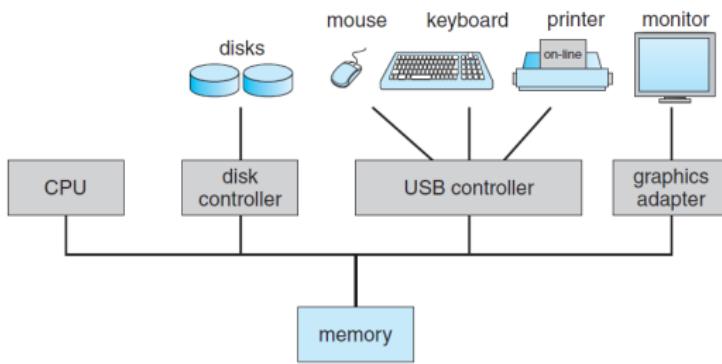
# Computer-System Operation

- ▶ One or more **CPUs**, and **device controllers** connect through common bus providing access to **shared memory**.



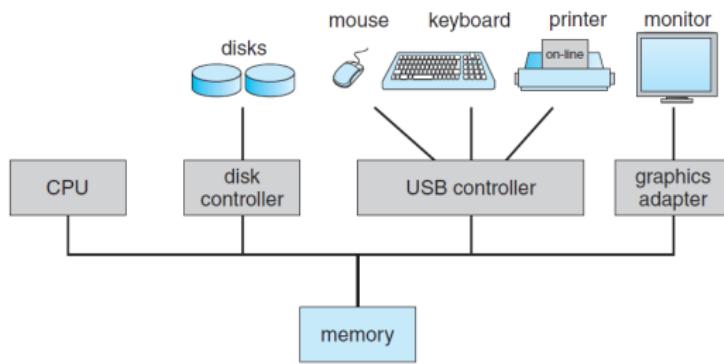
# Computer-System Operation

- ▶ One or more **CPUs**, and **device controllers** connect through common bus providing access to **shared memory**.
- ▶ The CPU and the device controllers can execute in **parallel**, **competing** for memory cycles.



# Computer-System Operation

- ▶ One or more **CPUs**, and **device controllers** connect through common bus providing access to **shared memory**.
- ▶ The CPU and the device controllers can execute in **parallel**, **competing** for memory cycles.
- ▶ Device controllers inform CPU that it is finished with the operation by causing an **interrupt**.



## Interrupt (1/3)

- ▶ **Hardware** may trigger an interrupt at any time by sending a **signal** to the CPU.

## Interrupt (1/3)

- ▶ **Hardware** may trigger an interrupt at any time by sending a **signal** to the CPU.
- ▶ **Software** may trigger an interrupt by executing a **special operation** called a **system call**.

## Interrupt (2/3)

- ▶ When the CPU is interrupted, it **stops** what it is doing and **immediately** transfers execution to an address where the **service routine** for the interrupt is located.

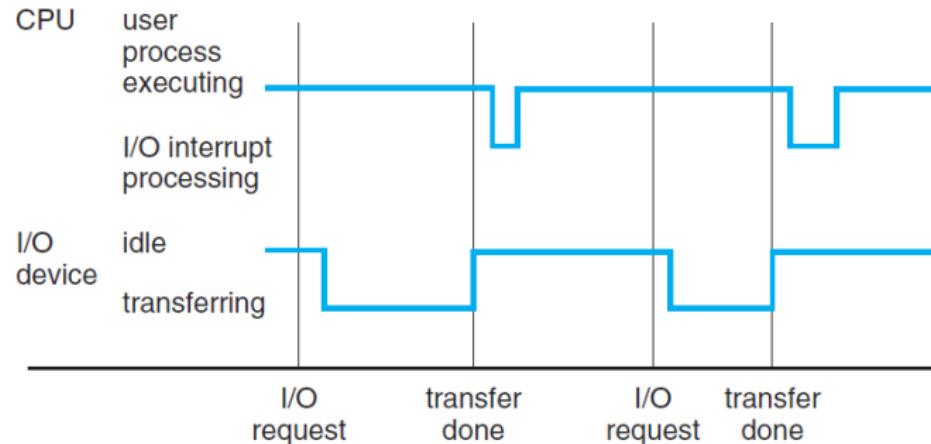
## Interrupt (2/3)

- ▶ When the CPU is interrupted, it **stops** what it is doing and **immediately** transfers execution to an address where the **service routine** for the interrupt is located.
- ▶ The CPU **resumes the interrupted computation**, when the **interrupt service routine** completes.

## Interrupt (2/3)

- ▶ When the CPU is interrupted, it **stops** what it is doing and **immediately** transfers execution to an address where the **service routine** for the interrupt is located.
- ▶ The CPU **resumes the interrupted computation**, when the **interrupt service routine** completes.
- ▶ The OS **preserves** the **state of the CPU** by storing registers and the program counter.

# Interrupt (3/3)



## ► Main memory (RAM)

- The CPU can **load instructions** only from memory.
- Typically **volatile**

## ► Main memory (RAM)

- The CPU can **load instructions** only from memory.
- Typically **volatile**

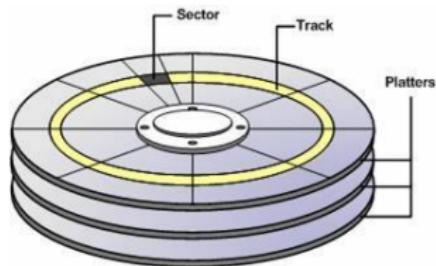
## ► Secondary storage

- **Extension** of main memory that provides large **nonvolatile** storage capacity.
- E.g., magnetic disk and SSD

## Storage Structure (2/2)

### ► Magnetic disks

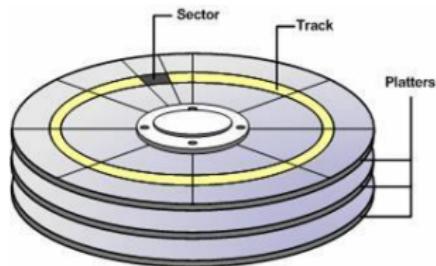
- Disk surface is logically divided into **tracks**, which are subdivided into **sectors**.



# Storage Structure (2/2)

## ► Magnetic disks

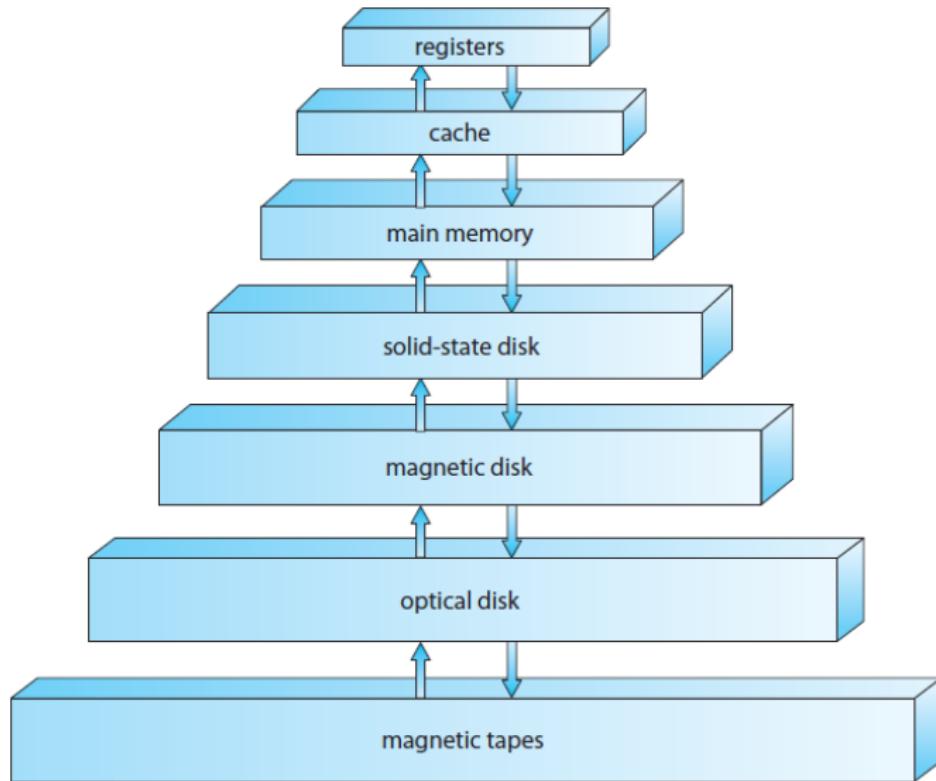
- Disk surface is logically divided into **tracks**, which are subdivided into **sectors**.



## ► Solid-state disks (SSD)

- **Faster** than hard disks, nonvolatile.
- Becoming more popular.

# Storage Hierarchy - Cost vs. Speed



# Summary

## ► OS history

- First generation: no OS
- Second generation: mainframes, batch programming
- Third generation: multiprogramming, multitasking
- Fourth generation: PCs

## ► OS history

- First generation: no OS
- Second generation: mainframes, batch programming
- Third generation: multiprogramming, multitasking
- Fourth generation: PCs

## ► Computer-system organization

- I/O devices
- Interrupt
- Storage

# Questions?

## Acknowledgements

Some slides were derived from Avi Silberschatz slides.