

# The Spark Big Data Analytics Platform

Amir H. Payberah

Swedish Institute of Computer Science

[amir@sics.se](mailto:amir@sics.se)

June 17, 2014





**"THAT'S your Ark for the Big Data  
flood? Noah, you will need a  
lot more storage space!"**

- ▶ Big Data refers to datasets and flows large enough that has outpaced our capability to store, process, analyze, and understand.



# Where Does Big Data Come From?

# Big Data Market Driving Factors

The number of web pages indexed by Google, which were around one million in 1998, have exceeded one trillion in 2008, and its expansion is accelerated by appearance of the social networks.\*



\* "Mining big data: current status, and forecast to the future" [Wei Fan et al., 2013]

# Big Data Market Driving Factors

The amount of **mobile data traffic** is expected to grow to **10.8 Exabyte** per month by **2016.\***



\* "Worldwide Big Data Technology and Services 2012-2015 Forecast" [Dan Vasset et al., 2013]

# Big Data Market Driving Factors

More than **65 billion devices** were connected to the Internet by **2010**, and this number will go up to **230 billion** by **2020**.\*



\* "The Internet of Things Is Coming" [John Mahoney et al., 2013]

# Big Data Market Driving Factors

Many companies are moving towards using **Cloud services** to access **Big Data analytical tools**.



# Big Data Market Driving Factors

Open source communities



# How To Process Big Data?



## Scale Up vs. Scale Out (1/2)

- ▶ Scale **up** or scale **vertically**: adding **resources** to a **single node** in a system.
- ▶ Scale **out** or scale **horizontally**: adding **more nodes** to a system.




## Scale Up vs. Scale Out (2/2)

- ▶ Scale **up**: more **expensive** than scaling out.
- ▶ Scale **out**: more challenging for **fault tolerance** and **software development**.




# Taxonomy of Parallel Architectures

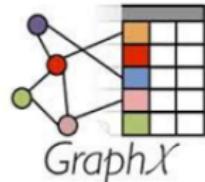


DeWitt, D. and Gray, J. "Parallel database systems: the future of high performance database systems". ACM Communications, 35(6), 85-98, 1992.

# Taxonomy of Parallel Architectures



DeWitt, D. and Gray, J. "Parallel database systems: the future of high performance database systems". ACM Communications, 35(6), 85-98, 1992.




Storm

**S4** distributed stream computing platform



# Big Data Analytics Stack



# Outline

- ▶ Introduction to **Scala**
- ▶ Data exploration using **Spark**
- ▶ Stream processing with **Spark Streaming**
- ▶ Graph analytics with **GraphX**



- ▶ **Scala**: scalable language
- ▶ A blend of object-oriented and functional programming
- ▶ Runs on the Java Virtual Machine
- ▶ Designed by Martin Odersky at EPFL



# Functional Programming Languages

- ▶ In a **restricted** sense: a language that does **not** have **mutable variables**, **assignments**, or **imperative control structures**.
- ▶ In a **wider** sense: it enables the construction of programs that **focus on functions**.

# Functional Programming Languages

- ▶ In a **restricted** sense: a language that does **not** have **mutable variables**, **assignments**, or **imperative control structures**.
- ▶ In a **wider** sense: it enables the construction of programs that **focus on functions**.
- ▶ **Functions** are **first-class citizens**:
  - **Defined anywhere** (including inside other functions).
  - **Passed as parameters** to functions and **returned as results**.
  - **Operators** to compose functions.

# Scala Variables

- ▶ **Values:** immutable
- ▶ **Variables:** mutable

```
var myVar: Int = 0  
val myVal: Int = 1
```

- ▶ Scala data types:
  - Boolean, Byte, Short, Char, Int, Long, Float, Double, String

## If ... Else

```
var x = 30;

if (x == 10) {
    println("Value of X is 10");
} else if (x == 20) {
    println("Value of X is 20");
} else {
    println("This is else statement");
}
```

# Loop

```
var a = 0
var b = 0
for (a <- 1 to 3; b <- 1 until 3) {
    println("Value of a: " + a + ", b: " + b )
}
```

```
// loop with collections
val numList = List(1, 2, 3, 4, 5, 6)
for (a <- numList) {
    println("Value of a: " + a)
}
```

# Functions

```
def functionName([list of parameters]): [return type] = {  
    function body  
    return [expr]  
}  
  
def addInt(a: Int, b: Int): Int = {  
    var sum: Int = 0  
    sum = a + b  
    sum  
}  
  
println("Returned Value: " + addInt(5, 7))
```

# Anonymous Functions

- ▶ Lightweight syntax for defining functions.

```
var mul = (x: Int, y: Int) => x * y
println(mul(3, 4))
```

# Higher-Order Functions



```
def apply(f: Int => String, v: Int) = f(v)

def layout(x: Int) = "[" + x.toString() + "]"

println(apply(layout, 10))
```

## Collections (1/2)

- ▶ **Array**: fixed-size sequential collection of elements of the same type

```
val t = Array("zero", "one", "two")
val b = t(0) // b = zero
```

## Collections (1/2)

- ▶ **Array**: fixed-size sequential collection of elements of the **same type**

```
val t = Array("zero", "one", "two")
val b = t(0) // b = zero
```

- ▶ **List**: sequential collection of elements of the **same type**

```
val t = List("zero", "one", "two")
val b = t(0) // b = zero
```

## Collections (1/2)

- ▶ **Array**: fixed-size sequential collection of elements of the **same type**

```
val t = Array("zero", "one", "two")
val b = t(0) // b = zero
```

- ▶ **List**: sequential collection of elements of the **same type**

```
val t = List("zero", "one", "two")
val b = t(0) // b = zero
```

- ▶ **Set**: sequential collection of elements of the **same type without duplicates**

```
val t = Set("zero", "one", "two")
val t.contains("zero")
```

## Collections (2/2)

- ▶ Map: collection of key/value pairs

```
val m = Map(1 -> "sics", 2 -> "kth")
val b = m(1) // b = sics
```

## Collections (2/2)

- ▶ Map: collection of key/value pairs

```
val m = Map(1 -> "sics", 2 -> "kth")
val b = m(1) // b = sics
```

- ▶ Tuple: A fixed number of items of different types together

```
val t = (1, "hello")
val b = t._1 // b = 1
val c = t._2 // c = hello
```

# Functional Combinators

- ▶ **map**: applies a function over each element in the list

```
val numbers = List(1, 2, 3, 4)
numbers.map(i => i * 2) // List(2, 4, 6, 8)
```

- ▶ **flatten**: it collapses one level of nested structure

```
List(List(1, 2), List(3, 4)).flatten // List(1, 2, 3, 4)
```

- ▶ **flatMap**: map + flatten

- ▶ **foreach**: it is like map but returns nothing

# Classes and Objects

```
class Calculator {  
    val brand: String = "HP"  
    def add(m: Int, n: Int): Int = m + n  
}  
  
val calc = new Calculator  
calc.add(1, 2)  
println(calc.brand)
```

# Classes and Objects

```
class Calculator {  
    val brand: String = "HP"  
    def add(m: Int, n: Int): Int = m + n  
}  
  
val calc = new Calculator  
calc.add(1, 2)  
println(calc.brand)
```

- ▶ A singleton is a class that can have only one instance.

```
object Test {  
    def main(args: Array[String]) { ... }  
}  
  
Test.main(null)
```

# Case Classes and Pattern Matching

- ▶ Case classes are used to store and match on the contents of a class.
- ▶ They are designed to be used with pattern matching.
- ▶ You can construct them without using new.

```
case class Calc(brand: String, model: String)

def calcType(calc: Calc) = calc match {
  case Calc("hp", "20B") => "financial"
  case Calc("hp", "48G") => "scientific"
  case Calc("hp", "30B") => "business"
  case _ => "Calculator of unknown type"
}

calcType(Calc("hp", "20B"))
```

# Simple Build Tool (SBT)

- ▶ An open source [build tool](#) for Scala and Java projects.
- ▶ Similar to Java's [Maven](#) or [Ant](#).
- ▶ It is written in [Scala](#).

# SBT - Hello World!

```
// make dir hello and edit Hello.scala
object Hello {
    def main(args: Array[String]) {
        println("Hello world.")
    }
}
```

```
$ cd hello
$ sbt compile run
```

# Common Commands

- ▶ `compile`: compiles the main sources.
- ▶ `run <argument>*`: run the main class.
- ▶ `package`: creates a jar file.
- ▶ `console`: starts the Scala interpreter.
- ▶ `clean`: deletes all generated files.
- ▶ `help <command>`: displays detailed help for the specified command.

# Create a Simple Project

- ▶ Create `project` directory.
- ▶ Create `src/main/scala` directory.
- ▶ Create `build.sbt` in the project root.

# build.sbt

- ▶ A list of Scala expressions, separated by blank lines.
- ▶ Located in the project's **base directory**.

```
$ cat build.sbt
name := "hello"

version := "1.0"

scalaVersion := "2.10.4"
```

# Add Dependencies

- ▶ Add in `build.sbt`.

- ▶ Module ID format:

```
"groupID" %% "artifact" % "version" % "configuration"
```

```
libraryDependencies += "org.apache.spark" %% "spark-core" % "1.0.0"

// multiple dependencies
libraryDependencies ++= Seq(
    "org.apache.spark" %% "spark-core" % "1.0.0",
    "org.apache.spark" %% "spark-streaming" % "1.0.0"
)
```

- ▶ sbt uses the standard Maven2 repository by default, but you can add more `resolvers`.

```
resolvers += "Akka Repository" at "http://repo.akka.io/releases/"
```

## Scala Hands-on Exercises (1/4)

- ▶ Declare a list of integers as a variable called `myNumbers`

## Scala Hands-on Exercises (1/4)

- ▶ Declare a list of integers as a variable called `myNumbers`

```
val myNumbers = List(1, 2, 5, 4, 7, 3)
```

## Scala Hands-on Exercises (1/4)

- ▶ Declare a list of integers as a variable called `myNumbers`

```
val myNumbers = List(1, 2, 5, 4, 7, 3)
```

- ▶ Declare a function, `pow`, that computes the second power of an Int

## Scala Hands-on Exercises (1/4)

- ▶ Declare a list of integers as a variable called `myNumbers`

```
val myNumbers = List(1, 2, 5, 4, 7, 3)
```

- ▶ Declare a function, `pow`, that computes the second power of an Int

```
def pow(a: Int): Int = a * a
```

## Scala Hands-on Exercises (2/4)

- ▶ Apply the function to `myNumbers` using the `map` function

## Scala Hands-on Exercises (2/4)

- ▶ Apply the function to `myNumbers` using the `map` function

```
myNumbers.map(x => pow(x))  
// or  
myNumbers.map(pow(_))  
// or  
myNumbers.map(pow)
```

## Scala Hands-on Exercises (2/4)

- ▶ Apply the function to `myNumbers` using the `map` function

```
myNumbers.map(x => pow(x))  
// or  
myNumbers.map(pow(_))  
// or  
myNumbers.map(pow)
```

- ▶ Write the `pow` function inline in a `map` call, using closure notation

## Scala Hands-on Exercises (2/4)

- ▶ Apply the function to `myNumbers` using the `map` function

```
myNumbers.map(x => pow(x))  
// or  
myNumbers.map(pow(_))  
// or  
myNumbers.map(pow)
```

- ▶ Write the `pow` function inline in a `map` call, using closure notation

```
myNumbers.map(x => x * x)
```

## Scala Hands-on Exercises (2/4)

- ▶ Apply the function to `myNumbers` using the `map` function

```
myNumbers.map(x => pow(x))  
// or  
myNumbers.map(pow(_))  
// or  
myNumbers.map(pow)
```

- ▶ Write the `pow` function inline in a `map` call, using closure notation

```
myNumbers.map(x => x * x)
```

- ▶ Iterate through `myNumbers` and print out its items

## Scala Hands-on Exercises (2/4)

- ▶ Apply the function to `myNumbers` using the `map` function

```
myNumbers.map(x => pow(x))  
// or  
myNumbers.map(pow(_))  
// or  
myNumbers.map(pow)
```

- ▶ Write the `pow` function inline in a `map` call, using closure notation

```
myNumbers.map(x => x * x)
```

- ▶ Iterate through `myNumbers` and print out its items

```
for (i <- myNumbers)  
  println(i)  
// or  
myNumbers.foreach(println)
```

## Scala Hands-on Exercises (3/4)

- ▶ Declare a list of pair of string and integers as a variable called `myList`

## Scala Hands-on Exercises (3/4)

- ▶ Declare a list of pair of string and integers as a variable called `myList`

```
val myList = List[(String, Int)](("a", 1), ("b", 2), ("c", 3))
```

## Scala Hands-on Exercises (3/4)

- ▶ Declare a list of pair of string and integers as a variable called `myList`

```
val myList = List[(String, Int)](("a", 1), ("b", 2), ("c", 3))
```

- ▶ Write an inline function to increment the integer values of the list `myList`

## Scala Hands-on Exercises (3/4)

- Declare a list of pair of string and integers as a variable called `myList`

```
val myList = List[(String, Int)](("a", 1), ("b", 2), ("c", 3))
```

- Write an inline function to increment the integer values of the list `myList`

```
val x = v.map { case (name, age) => age + 1 }
// or
val x = v.map(i => i._2 + 1)
// or
val x = v.map(_._2 + 1)
```

## Scala Hands-on Exercises (4/4)

- ▶ Do a word-count of a text file: create a Map with words as keys and counts of the number of occurrences of the word as values

## Scala Hands-on Exercises (4/4)

- ▶ Do a word-count of a text file: create a Map with words as keys and counts of the number of occurrences of the word as values
- ▶ You can load a text file as an array of lines as shown below:

```
import scala.io.Source  
val lines = Source.fromFile("/root/spark/README.md").getLines.toArray
```

## Scala Hands-on Exercises (4/4)

- ▶ Do a word-count of a text file: create a Map with words as keys and counts of the number of occurrences of the word as values
- ▶ You can load a text file as an array of lines as shown below:

```
import scala.io.Source  
val lines = Source.fromFile("/root/spark/README.md").getLines.toArray
```

- ▶ Then, instantiate a `HashMap[String, Int]` and use functional methods to populate it with word-counts

## Scala Hands-on Exercises (4/4)

- ▶ Do a word-count of a text file: create a Map with words as keys and counts of the number of occurrences of the word as values
- ▶ You can load a text file as an array of lines as shown below:

```
import scala.io.Source  
val lines = Source.fromFile("/root/spark/README.md").getLines.toArray
```

- ▶ Then, instantiate a `HashMap[String, Int]` and use functional methods to populate it with word-counts



```
val counts = new collection.mutable.HashMap[String, Int].withDefaultValue(0)  
lines.flatMap(_.split("""\W+""")).foreach(word => counts(word) += 1)  
counts.foreach(println)
```



# What is Spark?

- ▶ An efficient **distributed** general-purpose data analysis platform.
- ▶ Focusing on **ease** of programming and **high** performance.

# Spark Big Data Analytics Stack



# Motivation


- ▶ MapReduce programming model has not been designed for **complex** operations, e.g., data mining.
- ▶ Very **expensive**, i.e., always goes to disk and HDFS.

## Solution



- ▶ Extends MapReduce with **more** operators.
- ▶ Support for advanced **data flow graphs**.
- ▶ **In-memory** and **out-of-core** processing.




# Spark vs. Hadoop





# Spark vs. Hadoop



# Spark vs. Hadoop



# Spark vs. Hadoop



## Resilient Distributed Datasets (RDD) (1/2)


- ▶ A distributed memory abstraction.

## Resilient Distributed Datasets (RDD) (1/2)

- ▶ A **distributed memory** abstraction.
- ▶ **Immutable collections** of **objects** spread across a cluster.

## Resilient Distributed Datasets (RDD) (2/2)

- An **RDD** is divided into a number of **partitions**, which are **atomic** pieces of information.



- Partitions of an RDD can be stored on different **nodes** of a cluster.

# RDD Operators


- ▶ Higher-order functions: **transformations** and **actions**.
- ▶ **Transformations**: **lazy** operators that create **new** RDDs.
- ▶ **Actions**: launch a **computation** and return a **value** to the program or write data to the external storage.

# Transformations vs. Actions

Transformations	$map(f : T \Rightarrow U)$ : $RDD[T] \Rightarrow RDD[U]$ $filter(f : T \Rightarrow Bool)$ : $RDD[T] \Rightarrow RDD[T]$ $flatMap(f : T \Rightarrow Seq[U])$ : $RDD[T] \Rightarrow RDD[U]$ $sample(fraction : Float)$ : $RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) $groupByKey()$ : $RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ $reduceByKey(f : (V, V) \Rightarrow V)$ : $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $union()$ : $(RDD[T], RDD[T]) \Rightarrow RDD[T]$ $join()$ : $(RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ $cogroup()$ : $(RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ $crossProduct()$ : $(RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ $mapValues(f : V \Rightarrow W)$ : $RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) $sort(c : Comparator[K])$ : $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $partitionBy(p : Partitioner[K])$ : $RDD[(K, V)] \Rightarrow RDD[(K, V)]$
Actions	$count()$ : $RDD[T] \Rightarrow Long$ $collect()$ : $RDD[T] \Rightarrow Seq[T]$ $reduce(f : (T, T) \Rightarrow T)$ : $RDD[T] \Rightarrow T$ $lookup(k : K)$ : $RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs) $save(path : String)$ : Outputs RDD to a storage system, e.g., HDFS


## RDD Transformations - Map

- ▶ All pairs are **independently** processed.



## RDD Transformations - Map


- ▶ All pairs are **independently** processed.



```
// passing each element through a function.  
val nums = sc.parallelize(Array(1, 2, 3))  
val squares = nums.map(x => x * x) // {1, 4, 9}
```


## RDD Transformations - GroupBy

- ▶ Pairs with **identical key** are grouped.
- ▶ Groups are independently processed.



## RDD Transformations - GroupBy


- ▶ Pairs with **identical key** are grouped.
- ▶ Groups are independently processed.



```
val schools = sc.parallelize(Seq(("sics", 1), ("kth", 1), ("sics", 2)))  
  
schools.groupByKey()  
// {("sics", (1, 2)), ("kth", (1))}  
  
schools.reduceByKey((x, y) => x + y)  
// {("sics", 3), ("kth", 1)}
```


## RDD Transformations - Join

- ▶ Performs an **equi-join** on the key.
- ▶ Join candidates are independently processed.



## RDD Transformations - Join

- ▶ Performs an **equi-join** on the key.
- ▶ Join candidates are independently processed.



```
val list1 = sc.parallelize(Seq(("sics", "10"),
                             ("kth", "50"),
                             ("sics", "20")))

val list2 = sc.parallelize(Seq(("sics", "upsala"),
                             ("kth", "stockholm")))

list1.join(list2)
// ("sics", ("10", "upsala"))
// ("sics", ("20", "upsala"))
// ("kth", ("50", "stockholm"))
```

# Basic RDD Actions

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))
nums.collect() // Array(1, 2, 3)
```

# Basic RDD Actions

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))
nums.collect() // Array(1, 2, 3)
```

- ▶ Return an array with the first n elements of the RDD.

```
nums.take(2) // Array(1, 2)
```

# Basic RDD Actions

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))
nums.collect() // Array(1, 2, 3)
```

- ▶ Return an array with the first n elements of the RDD.

```
nums.take(2) // Array(1, 2)
```

- ▶ Return the number of elements in the RDD.

```
nums.count() // 3
```

# Basic RDD Actions

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))
nums.collect() // Array(1, 2, 3)
```

- ▶ Return an array with the first n elements of the RDD.

```
nums.take(2) // Array(1, 2)
```

- ▶ Return the number of elements in the RDD.

```
nums.count() // 3
```

- ▶ Aggregate the elements of the RDD using the given function.

```
nums.reduce((x, y) => x + y) // 6
```

# Creating RDDs

- ▶ Turn a collection into an RDD.

```
val a = sc.parallelize(Array(1, 2, 3))
```

- ▶ Load text file from local FS, HDFS, or S3.

```
val a = sc.textFile("file.txt")
val b = sc.textFile("directory/*.txt")
val c = sc.textFile("hdfs://namenode:9000/path/file")
```

# SparkContext

- ▶ Main entry point to Spark functionality.
- ▶ Available in `shell` as variable `sc`.
- ▶ In `standalone` programs, you should make your own.

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

val sc = new SparkContext(master, appName, [sparkHome], [jars])
```

## Spark Hands-on Exercises (1/3)

- ▶ Read data from the given file `hamlet` and create an RDD named `pagecounts`

## Spark Hands-on Exercises (1/3)

- ▶ Read data from the given file `hamlet` and create an RDD named `pagecounts`

```
val pagecounts = sc.textFile("hamlet")
```

## Spark Hands-on Exercises (1/3)

- ▶ Read data from the given file `hamlet` and create an RDD named `pagecounts`

```
val pagecounts = sc.textFile("hamlet")
```

- ▶ Get the first 10 lines of `hamlet`

## Spark Hands-on Exercises (1/3)

- ▶ Read data from the given file `hamlet` and create an RDD named `pagecounts`

```
val pagecounts = sc.textFile("hamlet")
```

- ▶ Get the first 10 lines of `hamlet`

```
pagecounts.take(10).foreach(println)
```

## Spark Hands-on Exercises (1/3)

- ▶ Read data from the given file `hamlet` and create an RDD named `pagecounts`

```
val pagecounts = sc.textFile("hamlet")
```

- ▶ Get the first 10 lines of `hamlet`

```
pagecounts.take(10).foreach(println)
```

- ▶ Count the total records in the data set `pagecounts`

## Spark Hands-on Exercises (1/3)

- ▶ Read data from the given file `hamlet` and create an RDD named `pagecounts`

```
val pagecounts = sc.textFile("hamlet")
```

- ▶ Get the first 10 lines of `hamlet`

```
pagecounts.take(10).foreach(println)
```

- ▶ Count the total records in the data set `pagecounts`

```
pagecounts.count
```

## Spark Hands-on Exercises (2/3)

- ▶ Filter the data set `pagecounts` and return the items that have the word `this`, and cache in the memory

## Spark Hands-on Exercises (2/3)

- ▶ Filter the data set `pagecounts` and return the items that have the word `this`, and cache in the memory

```
val linesWithThis = pagecounts.filter(line => line.contains("this")).cache  
\\ or  
val linesWithThis = pagecounts.filter(_.contains("this")).cache
```

## Spark Hands-on Exercises (2/3)

- ▶ Filter the data set `pagecounts` and return the items that have the word `this`, and cache in the memory

```
val linesWithThis = pagecounts.filter(line => line.contains("this")).cache  
\\ or  
val linesWithThis = pagecounts.filter(_.contains("this")).cache
```

- ▶ Find the lines with the most number of words.

## Spark Hands-on Exercises (2/3)

- ▶ Filter the data set `pagecounts` and return the items that have the word `this`, and cache in the memory

```
val linesWithThis = pagecounts.filter(line => line.contains("this")).cache  
\\ or  
val linesWithThis = pagecounts.filter(_.contains("this")).cache
```

- ▶ Find the lines with the most number of words.

```
linesWithThis.map(line => line.split(" ").size)  
.reduce((a, b) => if (a > b) a else b)
```

## Spark Hands-on Exercises (2/3)

- ▶ Filter the data set `pagecounts` and return the items that have the word `this`, and cache in the memory

```
val linesWithThis = pagecounts.filter(line => line.contains("this")).cache  
\\ or  
val linesWithThis = pagecounts.filter(_.contains("this")).cache
```

- ▶ Find the lines with the most number of words.

```
linesWithThis.map(line => line.split(" ").size)  
.reduce((a, b) => if (a > b) a else b)
```

- ▶ Count the total number of words

## Spark Hands-on Exercises (2/3)

- ▶ Filter the data set `pagecounts` and return the items that have the word `this`, and cache in the memory

```
val linesWithThis = pagecounts.filter(line => line.contains("this")).cache  
\\ or  
val linesWithThis = pagecounts.filter(_.contains("this")).cache
```

- ▶ Find the lines with the most number of words.

```
linesWithThis.map(line => line.split(" ").size)  
.reduce((a, b) => if (a > b) a else b)
```

- ▶ Count the total number of words

```
val wordCounts = linesWithThis.flatMap(line => line.split(" ")).count  
\\ or  
val wordCounts = linesWithThis.flatMap(_.split(" ")).count
```

## Spark Hands-on Exercises (3/3)

- ▶ Count the number of distinct words

## Spark Hands-on Exercises (3/3)

- ▶ Count the number of distinct words

```
val uniqueWordCounts = linesWithThis.flatMap(_.split(" ")).distinct.count
```

## Spark Hands-on Exercises (3/3)

- ▶ Count the number of distinct words

```
val uniqueWordCounts = linesWithThis.flatMap(_.split(" ")).distinct.count
```

- ▶ Count the number of each word

## Spark Hands-on Exercises (3/3)

- ▶ Count the number of distinct words

```
val uniqueWordCounts = linesWithThis.flatMap(_.split(" ")).distinct.count
```

- ▶ Count the number of each word

```
val eachWordCounts = linesWithThis.flatMap(_.split(" ")).  
    .map(word => (word, 1))  
    .reduceByKey((a, b) => a + b)
```

# Spark *Streaming*


# Motivation

- ▶ Many applications must process large **streams of live data** and provide results in **real-time**.
- ▶ Processing information as it **flows**, **without storing** them persistently.

- ▶ Many applications must process large **streams of live data** and provide results in **real-time**.
- ▶ Processing information as it **flows**, **without storing** them persistently.
- ▶ Traditional **DBMSs**:
  - **Store** and **index** data before processing it.
  - Process data only when **explicitly** asked by the users.
  - Both aspects **contrast** with our requirements.


# DBMS vs. DSMS (1/3)

- ▶ DBMS: persistent data where updates are relatively infrequent.
- ▶ DSMS: transient data that is continuously updated.



## DBMS vs. DSMS (2/3)

- ▶ DBMS: runs queries just **once** to return a complete answer.
- ▶ DSMS: executes **standing queries**, which run **continuously** and provide updated answers as new data arrives.



## DBMS vs. DSMS (3/3)

- ▶ Despite these differences, DSMSs resemble DBMSs: both process incoming data through a sequence of transformations based on SQL operators, e.g., selections, aggregates, joins.

# Spark Streaming

- ▶ Run a streaming computation as a **series** of very **small**, **deterministic batch** jobs.


# Spark Streaming

- ▶ Run a streaming computation as a **series** of very **small, deterministic batch** jobs.
  - **Chop up** the live stream into batches of **X** seconds.
  - Spark treats each batch of data as **RDDs** and processes them using **RDD operations**.
  - Finally, the processed results of the RDD operations are returned in **batches**.




# Spark Streaming API (1/4)

- ▶ **DStream**: sequence of RDDs representing a stream of data.
  - TCP sockets, Twitter, HDFS, Kafka, ...



# Spark Streaming API (1/4)

- ▶ **DStream**: sequence of RDDs representing a stream of data.
  - TCP sockets, Twitter, HDFS, Kafka, ...




- ▶ Initializing Spark streaming

```
val ssc = new StreamingContext(master, appName, batchDuration,  
[sparkHome], [jars])
```


# Spark Streaming API (2/4)

- ▶ **Transformations**: modify data from one DStream to a new DStream.
  - Standard RDD operations (**stateless** operations): map, join, ...




# Spark Streaming API (2/4)

- ▶ **Transformations**: modify data from one DStream to a new DStream.
  - Standard RDD operations (**stateless** operations): map, join, ...



- **Stateful** operations: group all the records from a sliding window of the past time intervals into one RDD: window, reduceByAndWindow, ...



**Window length:** the duration of the window.

**Slide interval:** the interval at which the operation is performed.

- ▶ **Output operations**: send data to external entity
  - saveAsHadoopFiles, foreach, print, ...

## Spark Streaming API (3/4)

- ▶ **Output operations**: send data to external entity
  - saveAsHadoopFiles, foreach, print, ...
- ▶ Attaching input sources

```
ssc.textFileStream(directory)
ssc.socketStream(hostname, port)
```

## Spark Streaming API (4/4)

- ▶ Stream + Batch: It can be used to apply any RDD operation that is not exposed in the DStream API.

```
val spamInfoRDD = sparkContext.hadoopFile(...)  
// join data stream with spam information to do data cleaning  
val cleanedDStream = inputDStream.transform(_.join(spamInfoRDD).filter(...))
```

## Spark Streaming API (4/4)

- ▶ Stream + Batch: It can be used to apply any RDD operation that is not exposed in the DStream API.

```
val spamInfoRDD = sparkContext.hadoopFile(...)  
// join data stream with spam information to do data cleaning  
val cleanedDStream = inputDStream.transform(_.join(spamInfoRDD).filter(...))
```

- ▶ Stream + Interactive: Interactive queries on stream state from the Spark interpreter

```
freqs.slice("21:00", "21:05").topK(10)
```

## Spark Streaming API (4/4)

- ▶ Stream + Batch: It can be used to apply any RDD operation that is not exposed in the DStream API.

```
val spamInfoRDD = sparkContext.hadoopFile(...)  
// join data stream with spam information to do data cleaning  
val cleanedDStream = inputDStream.transform(_.join(spamInfoRDD).filter(...))
```

- ▶ Stream + Interactive: Interactive queries on stream state from the Spark interpreter

```
freqs.slice("21:00", "21:05").topK(10)
```

- ▶ Starting/stopping the streaming computation


```
ssc.start()  
ssc.stop()  
ssc.awaitTermination()
```

## Example 1 (1/3)

- ▶ Get hash-tags from Twitter.

```
val ssc = new StreamingContext("local[2]", "Tweets", Seconds(1))  
val tweets = TwitterUtils.createStream(ssc, None)
```

DStream: a sequence of RDD representing a stream of data




## Example 1 (2/3)

- ▶ Get hash-tags from Twitter.

```
val ssc = new StreamingContext("local[2]", "Tweets", Seconds(1))
val tweets = TwitterUtils.createStream(ssc, None)
val hashTags = tweets.flatMap(status => getTags(status))
```


transformation: modify data in one DStream  
to create another DStream



## Example 1 (3/3)

- ▶ Get hash-tags from Twitter.


```
val ssc = new StreamingContext("local[2]", "Tweets", Seconds(1))
val tweets = TwitterUtils.createStream(ssc, None)
val hashTags = tweets.flatMap(status => getTags(status))
hashTags.saveAsHadoopFiles("hdfs://...")
```



## Example 2

- ▶ Count frequency of words received every second.


```
val ssc = new StreamingContext("local[2]", "NetworkWordCount", Seconds(1))
val lines = ssc.socketTextStream(ip, port)
val words = lines.flatMap(_.split(" "))
val ones = words.map(x => (x, 1))
val freqs = ones.reduceByKey(_ + _)
```



## Example 3

- ▶ Count frequency of words received in last minute.



```
val ssc = new StreamingContext("local[2]", "NetworkWordCount", Seconds(1))
val lines = ssc.socketTextStream(ip, port)
val words = lines.flatMap(_.split(" "))
val ones = words.map(x => (x, 1))
val freqs = ones.reduceByKey(_ + _)
val freqs_60s = freqs.window(Seconds(60), Second(1)).reduceByKey(_ + _)
```



## Example 3 - Simpler Model

- ▶ Count frequency of words received in last minute.


```
val ssc = new StreamingContext("local[2]", "NetworkWordCount", Seconds(1))
val lines = ssc.socketTextStream(ip, port)
val words = lines.flatMap(_.split(" "))
val ones = words.map(x => (x, 1))
val freqs_60s = ones.reduceByKeyAndWindow(_ + _, Seconds(60), Seconds(1))
```




## Example 3 - Incremental Window Operators

- ▶ Count frequency of words received in last minute.

```
// Associative only  
freqs_60s = ones.reduceByKeyAndWindow(_ + _, Seconds(60), Seconds(1))  
  
// Associative and invertible  
freqs_60s = ones.reduceByKeyAndWindow(_ + _, _ - _, Seconds(60), Seconds(1))
```



Associative only



Associative and invertible

# Spark Streaming Hands-on Exercises (1/2)

- ▶ Stream data through a TCP connection and port 9999

```
nc -lk 9999
```

# Spark Streaming Hands-on Exercises (1/2)

- ▶ Stream data through a TCP connection and port 9999

```
nc -lk 9999
```

- ▶ import the streaming libraries

# Spark Streaming Hands-on Exercises (1/2)

- ▶ Stream data through a TCP connection and port 9999

```
nc -lk 9999
```

- ▶ import the streaming libraries

```
import org.apache.spark.streaming.{Seconds, StreamingContext}
import org.apache.spark.streaming.StreamingContext._
```

# Spark Streaming Hands-on Exercises (1/2)

- ▶ Stream data through a TCP connection and port 9999

```
nc -lk 9999
```

- ▶ import the streaming libraries

```
import org.apache.spark.streaming.{Seconds, StreamingContext}  
import org.apache.spark.streaming.StreamingContext._
```

- ▶ Print out the incoming stream every five seconds at port 9999

# Spark Streaming Hands-on Exercises (1/2)

- ▶ Stream data through a TCP connection and port 9999

```
nc -lk 9999
```

- ▶ import the streaming libraries

```
import org.apache.spark.streaming.{Seconds, StreamingContext}
import org.apache.spark.streaming.StreamingContext._
```

- ▶ Print out the incoming stream every five seconds at port 9999

```
val ssc = new StreamingContext("local[2]", "NetworkWordCount", Seconds(5))
val lines = ssc.socketTextStream("127.0.0.1", 9999)
lines.print()
```

## Spark Streaming Hands-on Exercises (1/2)

- ▶ Count the number of each word in the incoming stream every five seconds at port 9999

# Spark Streaming Hands-on Exercises (1/2)

- ▶ Count the number of each word in the incoming stream every five seconds at port 9999

```
import org.apache.spark.streaming.{Seconds, StreamingContext}
import org.apache.spark.streaming.StreamingContext._

val ssc = new StreamingContext("local[2]", "NetworkWordCount", Seconds(1))
val lines = ssc.socketTextStream("127.0.0.1", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(x => (x, 1))
val wordCounts = pairs.reduceByKey(_ + _)
wordCounts.print()
```

## Spark Streaming Hands-on Exercises (2/2)

- ▶ Extend the code to generate word count over last 30 seconds of data, and repeat the computation every 10 seconds

## Spark Streaming Hands-on Exercises (2/2)

- ▶ Extend the code to generate word count over last 30 seconds of data, and repeat the computation every 10 seconds

```
import org.apache.spark.streaming.{Seconds, StreamingContext}
import org.apache.spark.streaming.StreamingContext._

val ssc = new StreamingContext("local[2]", "NetworkWordCount", Seconds(5))
val lines = ssc.socketTextStream("127.0.0.1", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))
val windowedWordCounts = pairs
    .reduceByKeyAndWindow(_ + _, _ - _, Seconds(30), Seconds(10))
windowedWordCounts.print()
wordCounts.print()
```

## Spark Streaming Hands-on Exercises - Twitter (1/7)

- ▶ Twitter credential setup: to access Twitter's sample tweet stream

# Spark Streaming Hands-on Exercises - Twitter (1/7)

- ▶ Twitter credential setup: to access Twitter's sample tweet stream
- ▶ Open the link: <https://apps.twitter.com/>

The screenshot shows a web browser interface for Twitter's Application Management. At the top, there is a blue header bar with the Twitter logo and the text "Application Management". Below this, a yellow status bar displays the message "Status" and "Your application has been successfully deleted.". The main content area is titled "Twitter Apps" and contains a message "You don't currently have any Twitter Apps." Below this message is a button labeled "Create New App", which is circled in red. At the bottom of the page, there is a footer bar with the Twitter logo, the word "Tweet", and the number "2,319".

# Spark Streaming Hands-on Exercises - Twitter (2/7)

 Application Management

## Create an application

### Application details

**Name**

SICS-test

Your application name. This is used to attribute the source of a tweet and in user-facing authorization screens. 32 characters max.

**Description \***

App description

Your application description, which will be shown in user-facing authorization screens. Between 10 and 200 characters max.

**Website \***

http://www.sics.se

Your application's publicly accessible home page, where users can go to download, make use of, or find out more information about your application. This fully-qualified URL is used in the source attribution for tweets created by your application and will be shown in user-facing authorization screens.  
(If you don't have a URL yet, just put a placeholder here but remember to change it later.)

### Callback URL

Where should we return after successfully authenticating? OAuth 1.0a applications should explicitly specify their oauth\_callback URL on the request token step, regardless of the value given here. To restrict your application from using callbacks, leave this field blank.

sell, rent, lease, sublicense, redistribute, or syndicate access to the Twitter API or Twitter Content to any third party without prior written approval from Twitter.

Yes, I agree

**Create your Twitter application**

# Spark Streaming Hands-on Exercises - Twitter (3/7)

Application Management

## SICS-test

Test OAuth

Details   Settings   **API Keys**   Permissions

### Application settings

Keep the "API secret" a secret. This key should never be human-readable in your application.

API key	YNCP38MpSIUJfgCpWdhjCXy
API secret	4nQy5b7MOTpUoJ7BSW3KB75Amy2W0oCh3LMEFXCE85cRy4leZu
Access level	Read-only (modify app permissions)
Owner	payberah
Owner ID	42611668

### Application actions

Regenerate API keys   Change App Permissions

Your access token

You haven't authorized this application for your own account yet.

By creating your access token here, you will have everything you need to make API calls right away. The access token generated will be assigned your application's current permission level.

### Token actions

Create my access token

# Spark Streaming Hands-on Exercises - Twitter (4/7)

 Application Management

Status  
Your application access token has been successfully generated. It may take a moment for changes you've made to reflect.  
[Refresh](#) if your changes are not yet indicated.

## SICS-test

[Details](#) [Settings](#) [API Keys](#) [Permissions](#) [Test OAuth](#) (circled in red)

### Application settings

Keep the "API secret" a secret. This key should never be human-readable in your application.

API key	YNCP38MpSIUJfgCpWdhjCXy
API secret	4nQy5b7MOTpUoJ7BSW3KB75Amy2W0oCh3LMEFXCE85cRy4leZu
Access level	Read-only ( <a href="#">modify app permissions</a> )
Owner	payberah
Owner ID	42611668

# Spark Streaming Hands-on Exercises - Twitter (5/7)

The screenshot shows the Twitter Developers OAuth Settings page for a application named "SICS-test". The page has a header with links for Developers, API Health, Blog, Discussions, Documentation, and a search bar. Below the header, there's a breadcrumb navigation: Home → My applications. The main content area is titled "OAuth Settings" and contains four input fields for OAuth credentials, each with its value redacted:

- Consumer key:** [REDACTED]
- Consumer secret:** [REDACTED]  
Remember this should not be shared.
- Access token:** [REDACTED]
- Access token secret:** [REDACTED]  
Remember this should not be shared.

A "Manage" button is located in the top right corner of the main content area.

## Spark Streaming Hands-on Exercises - Twitter (6/7)

- ▶ Create an `StreamingContext` for a batch duration of 5 seconds and use this context to create a stream of tweets

## Spark Streaming Hands-on Exercises - Twitter (6/7)

- ▶ Create an `StreamingContext` for a batch duration of 5 seconds and use this context to create a stream of tweets

```
import org.apache.spark.streaming.twitter._

val ssc = new StreamingContext("local[2]", "Tweets", Seconds(5))
val tweets = TwitterUtils.createStream(ssc, None)
```

## Spark Streaming Hands-on Exercises - Twitter (6/7)

- ▶ Create an `StreamingContext` for a batch duration of 5 seconds and use this context to create a stream of tweets

```
import org.apache.spark.streaming.twitter._

val ssc = new StreamingContext("local[2]", "Tweets", Seconds(5))
val tweets = TwitterUtils.createStream(ssc, None)
```

- ▶ Print the status text of some of the tweets

# Spark Streaming Hands-on Exercises - Twitter (6/7)

- ▶ Create an `StreamingContext` for a batch duration of 5 seconds and use this context to create a stream of tweets

```
import org.apache.spark.streaming.twitter._

val ssc = new StreamingContext("local[2]", "Tweets", Seconds(5))
val tweets = TwitterUtils.createStream(ssc, None)
```

- ▶ Print the status text of some of the tweets

```
val statuses = tweets.map(status => status.getText())
statuses.print()
```

## Spark Streaming Hands-on Exercises - Twitter (6/7)

- ▶ Create an `StreamingContext` for a batch duration of 5 seconds and use this context to create a stream of tweets

```
import org.apache.spark.streaming.twitter._

val ssc = new StreamingContext("local[2]", "Tweets", Seconds(5))
val tweets = TwitterUtils.createStream(ssc, None)
```

- ▶ Print the status text of some of the tweets

```
val statuses = tweets.map(status => status.getText())
statuses.print()
```

- ▶ Get the stream of hashtags from the stream of tweets

# Spark Streaming Hands-on Exercises - Twitter (6/7)

- ▶ Create an `StreamingContext` for a batch duration of 5 seconds and use this context to create a stream of tweets

```
import org.apache.spark.streaming.twitter._

val ssc = new StreamingContext("local[2]", "Tweets", Seconds(5))
val tweets = TwitterUtils.createStream(ssc, None)
```

- ▶ Print the status text of some of the tweets

```
val statuses = tweets.map(status => status.getText())
statuses.print()
```

- ▶ Get the stream of hashtags from the stream of tweets

```
val words = statuses.flatMap(status => status.split(" "))
val hashtags = words.filter(word => word.startsWith("#"))
hashtags.print()
```

## Spark Streaming Hands-on Exercises - Twitter (7/7)

- ▶ Set a path for periodic checkpointing of the intermediate data, and then count the hashtags over a one minute window

# Spark Streaming Hands-on Exercises - Twitter (7/7)

- ▶ Set a path for periodic checkpointing of the intermediate data, and then count the hashtags over a one minute window

```
ssc.checkpoint("/home/sics/temp")
val counts = hashtags.map(tag => (tag, 1))
    .reduceByKeyAndWindow(_ + _, _ - _, Seconds(60), Seconds(5))
counts.print()
```

# Spark Streaming Hands-on Exercises - Twitter (7/7)

- ▶ Set a path for periodic checkpointing of the intermediate data, and then count the hashtags over a one minute window

```
ssc.checkpoint("/home/sics/temp")
val counts = hashtags.map(tag => (tag, 1))
    .reduceByKeyAndWindow(_ + _, _ - _, Seconds(60), Seconds(5))
counts.print()
```

- ▶ Find the top 10 hashtags based on their counts


# Spark Streaming Hands-on Exercises - Twitter (7/7)

- ▶ Set a path for periodic checkpointing of the intermediate data, and then count the hashtags over a one minute window

```
ssc.checkpoint("/home/sics/temp")
val counts = hashtags.map(tag => (tag, 1))
    .reduceByKeyAndWindow(_ + _, _ - _, Seconds(60), Seconds(5))
counts.print()
```

- ▶ Find the top 10 hashtags based on their counts

```
val sortedCounts = counts.map { case (tag, count) => (count, tag) }
    .transform(rdd => rdd.sortByKey(false))
sortedCounts.foreachRDD(rdd =>
    println("\nTop 10 hashtags:\n" + rdd.take(10).mkString("\n")))
```






# Introduction

- ▶ **Graphs** provide a flexible abstraction for describing relationships between **discrete objects**.
- ▶ Many problems can be modeled by graphs and solved with appropriate **graph algorithms**.

# Large Graph




# Large-Scale Graph Processing

- ▶ Large graphs need **large-scale processing**.
- ▶ A large graph either **cannot fit into memory** of single computer or it fits with huge cost.


# Question

Can we use platforms like MapReduce or Spark, which are based on data-parallel model, for large-scale graph proceeding?



# Data-Parallel Model for Large-Scale Graph Processing

- ▶ The platforms that have worked well for developing **parallel applications** are not necessarily effective for **large-scale graph** problems.
- ▶ Why?



# Graph Algorithms Characteristics (1/2)

## ► Unstructured problems

- Difficult to extract parallelism based on partitioning of the data: the irregular structure of graphs.
- Limited scalability: unbalanced computational loads resulting from poorly partitioned data.

# Graph Algorithms Characteristics (1/2)

## ► Unstructured problems

- Difficult to extract parallelism based on partitioning of the data: the irregular structure of graphs.
- Limited scalability: unbalanced computational loads resulting from poorly partitioned data.

## ► Data-driven computations

- Difficult to express parallelism based on partitioning of computation: the structure of computations in the algorithm is not known a priori.
- The computations are dictated by nodes and links of the graph.

## Graph Algorithms Characteristics (2/2)

- ▶ Poor data locality
  - The computations and data access patterns do not have much locality: the irregular structure of graphs.

## Graph Algorithms Characteristics (2/2)


- ▶ Poor data locality
  - The computations and data access patterns do not have much locality: the irregular structure of graphs.
- ▶ High data access to computation ratio
  - Graph algorithms are often based on exploring the structure of a graph to perform computations on the graph data.
  - Runtime can be dominated by waiting memory fetches: low locality.

# Proposed Solution

Graph-Parallel Processing

# Proposed Solution


## Graph-Parallel Processing



- ▶ Computation typically depends on the **neighbors**.


# Graph-Parallel Processing

- ▶ Restricts the **types of computation**.
- ▶ New techniques to **partition and distribute graphs**.
- ▶ Exploit graph structure.
- ▶ Executes graph algorithms orders-of-magnitude faster than more general **data-parallel** systems.



# Data-Parallel vs. Graph-Parallel Computation


Data-Parallel



Graph-Parallel




Property Graph



# Data-Parallel vs. Graph-Parallel Computation

- ▶ Data-parallel computation
  - Record-centric view of data.
  - Parallelism: processing **independent** data on separate resources.
  
- ▶ Graph-parallel computation
  - Vertex-centric view of graphs.
  - Parallelism: partitioning graph (**dependent**) data across processing resources, and **resolving dependencies (along edges)** through **iterative** computation and communication.



# Graph-Parallel Computation Frameworks




# Data-Parallel vs. Graph-Parallel Computation


- ▶ Graph-parallel computation: **restricting** the types of computation to achieve **performance**.

# Data-Parallel vs. Graph-Parallel Computation

- ▶ Graph-parallel computation: **restricting** the types of computation to achieve **performance**.
- ▶ **But**, the same restrictions make it **difficult** and **inefficient** to express many stages in a typical graph-analytics **pipeline**.




# Data-Parallel and Graph-Parallel Pipeline





- ▶ **Moving** between **table** and **graph** views of the **same physical data**.
- ▶ **Inefficient**: extensive **data movement** and **duplication** across the network and file system.

# GraphX vs. Data-Parallel/Graph-Parallel Systems



# GraphX vs. Data-Parallel/Graph-Parallel Systems




# GraphX

- ▶ New API that blurs the distinction between Tables and Graphs.
- ▶ New system that unifies Data-Parallel and Graph-Parallel systems.
- ▶ It is implemented on top of Spark.

# Unifying Data-Parallel and Graph-Parallel Analytics


- ▶ Tables and Graphs are composable views of the same physical data.
- ▶ Each view has its own operators that exploit the semantics of the view to achieve efficient execution.



# Data Model

- ▶ **Property Graph:** represented using two Spark RDDs:
  - Edge collection: VertexRDD
  - Vertex collection: EdgeRDD


```
// VD: the type of the vertex attribute
// ED: the type of the edge attribute
class Graph[VD, ED] {
    val vertices: VertexRDD[VD]
    val edges: EdgeRDD[ED, VD]
}
```




# Primitive Data Types

```
// Vertex collection
class VertexRDD[VD] extends RDD[(VertexId, VD)]  
  
// Edge collection
class EdgeRDD[ED] extends RDD[Edge[ED]]
case class Edge[ED, VD](srcId: VertexId = 0, dstId: VertexId = 0,
                       attr: ED = null.asInstanceOf[ED])  
  
// Edge Triple
class EdgeTriplet[VD, ED] extends Edge[ED]
```

- ▶ **EdgeTriplet** represents an **edge** along with the **vertex attributes** of its **neighboring** vertices.



# Example (1/3)



Vertex Table

Id	Property (V)
3	(rxin, student)
7	(jgonzal, postdoc)
5	(franklin, professor)
2	(istoica, professor)

Edge Table

SrcId	DstId	Property (E)
3	7	Collaborator
5	3	Advisor
2	5	Colleague
5	7	PI

## Example (2/3)

```
val sc: SparkContext

// Create an RDD for the vertices
val users: RDD[(Long, (String, String))] = sc.parallelize(
  Array((3L, ("rxin", "student")), (7L, ("jgonzal", "postdoc")),
    (5L, ("franklin", "prof")), (2L, ("istoica", "prof"))))

// Create an RDD for edges
val relationships: RDD[Edge[String]] = sc.parallelize(
  Array(Edge(3L, 7L, "collab"), Edge(5L, 3L, "advisor"),
    Edge(2L, 5L, "colleague"), Edge(5L, 7L, "pi")))

// Define a default user in case there are relationship with missing user
val defaultUser = ("John Doe", "Missing")

// Build the initial Graph
val userGraph: Graph[(String, String), String] =
  Graph(users, relationships, defaultUser)
```

## Example (3/3)

```
// Constructed from above
val userGraph: Graph[(String, String), String]

// Count all users which are postdocs
userGraph.vertices.filter { case (id, (name, pos)) => pos == "postdoc" }.count

// Count all the edges where src > dst
userGraph.edges.filter(e => e.srcId > e.dstId).count

// Use the triplets view to create an RDD of facts
val facts: RDD[String] = graph.triplets.map(triplet =>
  triplet.srcAttr._1 + " is the " +
  triplet.attr + " of " + triplet.dstAttr._1)

// Remove missing vertices as well as the edges to connected to them
val validGraph = graph.subgraph(vpred = (id, attr) => attr._2 != "Missing")

facts.collect.foreach(println)
```

## Property Operators (1/2)

```
class Graph[VD, ED] {  
    def mapVertices[VD2](map: (VertexId, VD) => VD2): Graph[VD2, ED]  
  
    def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]  
  
    def mapTriplets[ED2](map: EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]  
}
```

- ▶ They yield **new graphs** with the vertex or edge properties modified by the map function.
- ▶ The graph **structure** is **unaffected**.

## Property Operators (2/2)

```
val newGraph = graph.mapVertices((id, attr) => mapUdf(id, attr))
```

```
val newVertices = graph.vertices.map((id, attr) => (id, mapUdf(id, attr)))
val newGraph = Graph(newVertices, graph.edges)
```


- ▶ Both are logically equivalent, but the second one **does not preserve** the structural indices and would not benefit from the GraphX system **optimizations**.


# Map Reduce Triplets

- ▶ Map-Reduce for each vertex

mapF( ) → 

mapF( ) → 

reduceF( ) → 




# Map Reduce Triplets

- ▶ Map-Reduce for each vertex

mapF(  $\rightarrow$

mapF(  $\rightarrow$

reduceF(  $\rightarrow$



```
// what is the age of the oldest follower for each user?  
val oldestFollowerAge = graph.mapReduceTriplets(  
    e => Iterator((e.dstAttr, e.srcAttr)), // Map  
    (a, b) => max(a, b) // Reduce  
).vertices
```

# Structural Operators

```
class Graph[VD, ED] {  
    // returns a new graph with all the edge directions reversed  
    def reverse: Graph[VD, ED]  
  
    // returns the graph containing only the vertices and edges that satisfy  
    // the vertex predicate  
    def subgraph(epred: EdgeTriplet[VD,ED] => Boolean,  
                vpred: (VertexId, VD) => Boolean): Graph[VD, ED]  
  
    // a subgraph by returning a graph that contains the vertices and edges  
    // that are also found in the input graph  
    def mask[VD2, ED2](other: Graph[VD2, ED2]): Graph[VD, ED]  
}
```

# Structural Operators Example

```
// Build the initial Graph
val graph = Graph(users, relationships, defaultUser)

// Run Connected Components
val ccGraph = graph.connectedComponents()

// Remove missing vertices as well as the edges to connected to them
val validGraph = graph.subgraph(vpred = (id, attr) => attr._2 != "Missing")


// Restrict the answer to the valid subgraph
val validCCGraph = ccGraph.mask(validGraph)
```

# Join Operators

- ▶ To join data from external collections (RDDs) with graphs.

```
class Graph[VD, ED] {  
    // joins the vertices with the input RDD and returns a new graph  
    // by applying the map function to the result of the joined vertices  
    def joinVertices[U](table: RDD[(VertexId, U)])  
        (map: (VertexId, VD, U) => VD): Graph[VD, ED]  
  
    // similarly to joinVertices, but the map function is applied to  
    // all vertices and can change the vertex property type  
    def outerJoinVertices[U, VD2](table: RDD[(VertexId, U)])  
        (map: (VertexId, VD, Option[U]) => VD2): Graph[VD2, ED]  
}
```

# GraphX Hands-on Exercises (1/7)



## GraphX Hands-on Exercises (2/7)

- ▶ import the streaming libraries

## GraphX Hands-on Exercises (2/7)

- ▶ import the streaming libraries

```
import org.apache.spark.graphx._  
import org.apache.spark.rdd.RDD
```

## GraphX Hands-on Exercises (2/7)

- ▶ import the streaming libraries

```
import org.apache.spark.graphx._  
import org.apache.spark.rdd.RDD
```

- ▶ Build the property graph shown in the last page

## GraphX Hands-on Exercises (2/7)

- ▶ import the streaming libraries

```
import org.apache.spark.graphx._  
import org.apache.spark.rdd.RDD
```

- ▶ Build the property graph shown in the last page

```
val vertexArray = Array(  
    (1L, ("Alice", 28)), (2L, ("Bob", 27)), (3L, ("Charlie", 65)),  
    (4L, ("David", 42)), (5L, ("Ed", 55)), (6L, ("Fran", 50)))  
  
val edgeArray = Array(  
    Edge(2L, 1L, 7), Edge(2L, 4L, 2), Edge(3L, 2L, 4),  
    Edge(3L, 6L, 3), Edge(4L, 1L, 1), Edge(5L, 2L, 2),  
    Edge(5L, 3L, 8), Edge(5L, 6L, 3))  
  
val vertexRDD: RDD[(Long, (String, Int))] = sc.parallelize(vertexArray)  
val edgeRDD: RDD[Edge[Int]] = sc.parallelize(edgeArray)  
val graph: Graph[(String, Int), Int] = Graph(vertexRDD, edgeRDD)
```

## GraphX Hands-on Exercises (3/7)

- ▶ Display the name of the users older than 30 years old

## GraphX Hands-on Exercises (3/7)

- ▶ Display the name of the users older than 30 years old

```
graph.vertices.filter { case (id, (name, age)) => age > 30 }.foreach {  
    case (id, (name, age)) => println(s"$name is $age")  
}
```

## GraphX Hands-on Exercises (3/7)

- ▶ Display the name of the users older than 30 years old

```
graph.vertices.filter { case (id, (name, age)) => age > 30 }.foreach {  
    case (id, (name, age)) => println(s"$name is $age")  
}
```

- ▶ Display who follows who (through the edges direction).

```
/**  
 * Triplet has the following Fields:  
 *   triplet.srcAttr: (String, Int)  
 *   triplet.dstAttr: (String, Int)  
 *   triplet.attr: Int  
 *   triplet.srcId: VertexId  
 *   triplet.dstId: VertexId  
 */
```

## GraphX Hands-on Exercises (3/7)

- ▶ Display the name of the users older than 30 years old

```
graph.vertices.filter { case (id, (name, age)) => age > 30 }.foreach {  
    case (id, (name, age)) => println(s"$name is $age")  
}
```

- ▶ Display who follows who (through the edges direction).

```
/**  
 * Triplet has the following Fields:  
 *   triplet.srcAttr: (String, Int)  
 *   triplet.dstAttr: (String, Int)  
 *   triplet.attr: Int  
 *   triplet.srcId: VertexId  
 *   triplet.dstId: VertexId  
 */
```

```
graph.triplets.foreach(t =>  
    println(s"${t.srcAttr._1} follows ${t.dstAttr._1}"))
```

## GraphX Hands-on Exercises (4/7)

- ▶ Compute the total age of followers of each user and print them out

## GraphX Hands-on Exercises (4/7)

- ▶ Compute the total age of followers of each user and print them out

```
val followers: VertexRDD[Int] = graph.mapReduceTriplets[Int] (
    triplet => Iterator(...), // map
    (a, b) => ... // reduce
)
```

# GraphX Hands-on Exercises (4/7)

- ▶ Compute the total age of followers of each user and print them out

```
val followers: VertexRDD[Int] = graph.mapReduceTriplets[Int] (   
    triplet => Iterator(...), // map  
    (a, b) => ... // reduce  
)
```

```
val followers: VertexRDD[Int] = graph.mapReduceTriplets[Int] (   
    triplet => Iterator((triplet.dstId, triplet.srcAttr._2)),  
    (a, b) => a + b)  
  
followers.collect.foreach(print)
```

## GraphX Hands-on Exercises (5/7)

- ▶ Compute the average age of followers of each user and print them out

## GraphX Hands-on Exercises (5/7)

- ▶ Compute the average age of followers of each user and print them out

```
val followers: VertexRDD[(Int, Double)] = graph
  .mapReduceTriplets[(Int, Double)](
    triplet => Iterator(...), // map
    (a, b) => (...) // reduce
  )

val avgAgeOfFollowers: VertexRDD[Double] = followers.mapValues(...)
```

# GraphX Hands-on Exercises (5/7)

- ▶ Compute the average age of followers of each user and print them out

```
val followers: VertexRDD[(Int, Double)] = graph
  .mapReduceTriplets[(Int, Double)](
    triplet => Iterator(...), // map
    (a, b) => (...) // reduce
  )

val avgAgeOfFollowers: VertexRDD[Double] = followers.mapValues(...)
```

```
val followers: VertexRDD[(Int, Double)] = graph
  .mapReduceTriplets[(Int, Double)](
    triplet => Iterator((triplet.dstId, (1, triplet.srcAttr._2))),
    (a, b) => (a._1 + b._1, a._2 + b._2))

val avgAgeOfFollowers: VertexRDD[Double] =
  followers.mapValues((id, value) => value match {
    case (count, totalAge) => totalAge / count
  })

avgAgeOfFollowers.collect.foreach(print)
```

## GraphX Hands-on Exercises (6/7)

- ▶ Make a subgraph of the users that are 30 or older

## GraphX Hands-on Exercises (6/7)

- ▶ Make a subgraph of the users that are 30 or older

```
val olderGraph = graph.subgraph(vpred = ...)
```

# GraphX Hands-on Exercises (6/7)

- ▶ Make a subgraph of the users that are 30 or older

```
val olderGraph = graph.subgraph(vpred = ...)
```

```
val olderGraph = graph.subgraph(vpred = (id, u) => u._2 >= 30)
```

## GraphX Hands-on Exercises (7/7)

- ▶ Compute the connected components and display the component id of each user in oldGraph

## GraphX Hands-on Exercises (7/7)

- ▶ Compute the connected components and display the component id of each user in oldGraph

```
val cc = olderGraph...

olderGraph.vertices.leftJoin(cc.vertices) {
...
}.foreach{...}
```

# GraphX Hands-on Exercises (7/7)

- ▶ Compute the connected components and display the component id of each user in oldGraph

```
val cc = olderGraph...

olderGraph.vertices.leftJoin(cc.vertices) {
...
}.foreach{...}
```

```
val cc = olderGraph.connectedComponents

olderGraph.vertices.leftJoin(cc.vertices) {
  case (id, u, comp) => s"${u._1} is in component ${comp.get}"
}.foreach{ case (id, str) => println(str) }
```

# Questions?